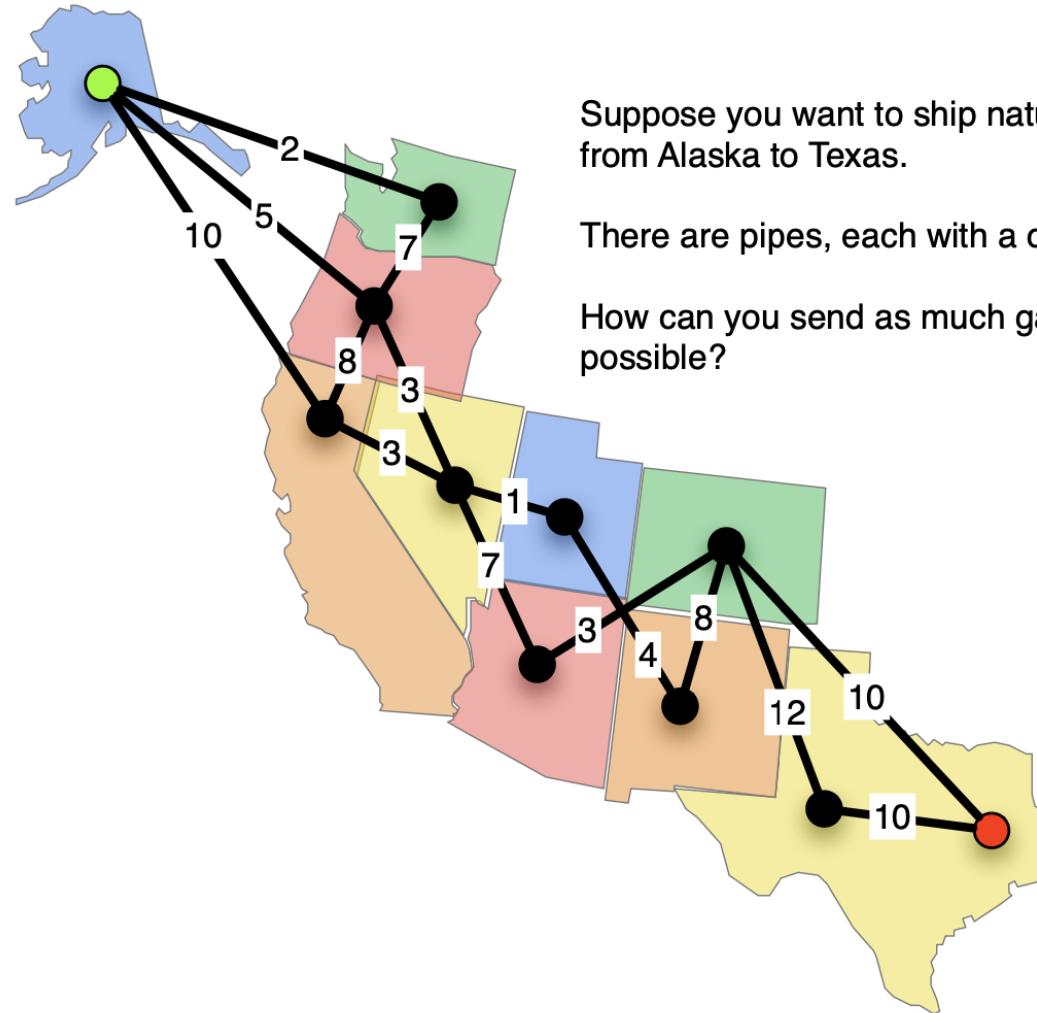


Design and Analysis of Algorithms

Jing Tang | DSAA 2043 Fall 2024

Network Flow

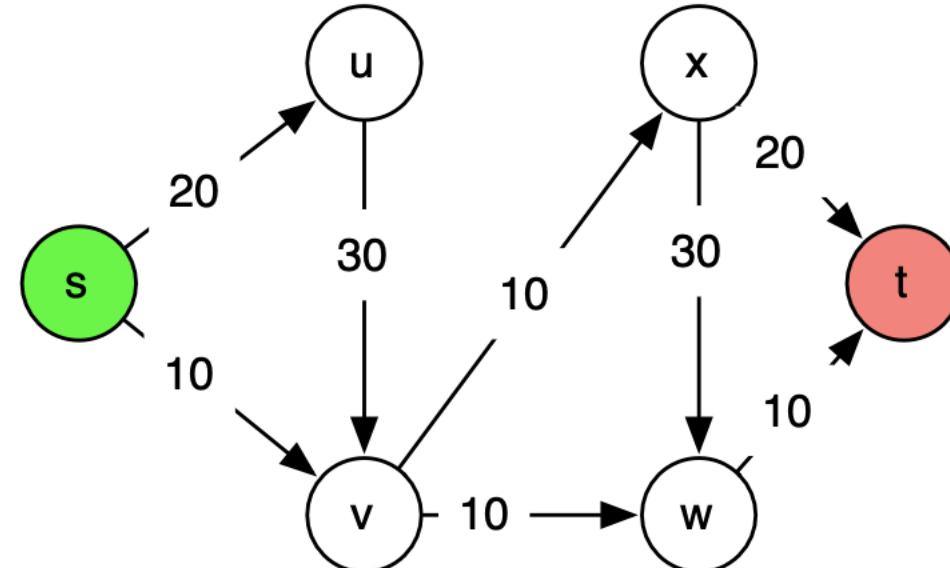
Network Flow Problem e.g.



Flow Network

A flow network is a connected, directed graph $G = (V, E)$.

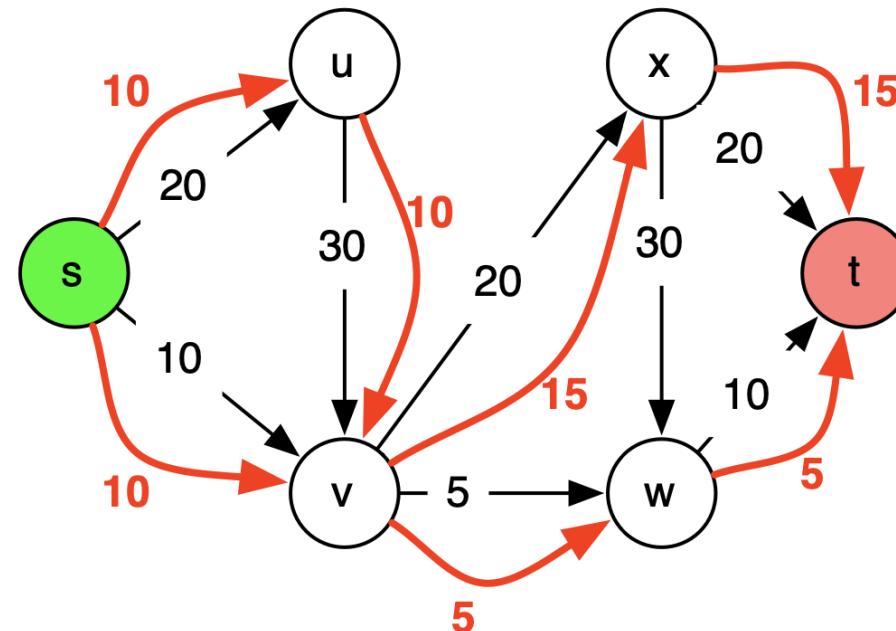
- Each edge e has a non-negative, integer capacity c_e .
- A single source $s \in V$.
- A single sink $t \in V$.
- No edge enters the source and no edge leaves the sink.



Assumptions

- To repeat, we make these assumptions about the network:
 - Capacities are integers.
 - Every node has one edge adjacent to it.
 - No edge enters the source and no edge leaves the sink.
- These assumptions can all be removed.

- Def. An s - t flow is a function $f: E \rightarrow \mathbb{R}_{\geq 0}$ that assigns a real number to each edge.
- Intuitively, $f(e)$ is the amount of material carried on the edge e .



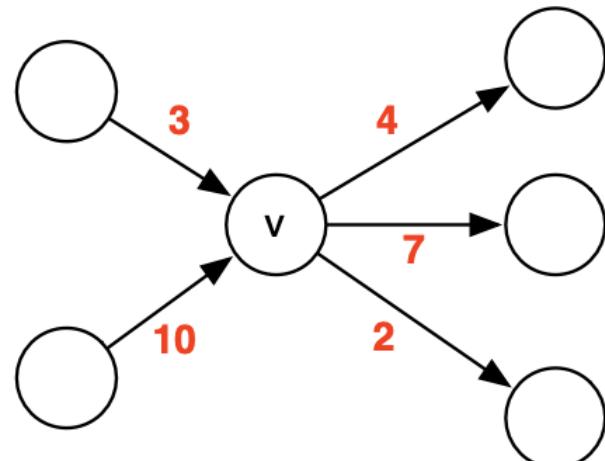
Flow Constraints

Constraints on f :

- $0 \leq f(e) \leq c_e$ for each edge e . (capacity constraints)
- For each node v except s and t , we have:

$$\sum_{e \text{ into } v} f(e) = \sum_{e \text{ leaving } v} f(e)$$

(balance constraints: whatever flows in, must flow out).



Notation

- The value of flow f is:

$$v(f) = \sum_{e \text{ out of } s} f(e)$$

This is the amount of material that s is able to send out.

- Notation:

$$- f^{\text{in}}(v) = \sum_{e \text{ into } v} f(e)$$

$$- f^{\text{out}}(v) = \sum_{e \text{ leaving } v} f(e)$$

- Balance constraints becomes: $f^{\text{in}}(v) = f^{\text{out}}(v)$ for all $v \notin \{s,t\}$

Maximum Flow Problem

- Definition (Value)

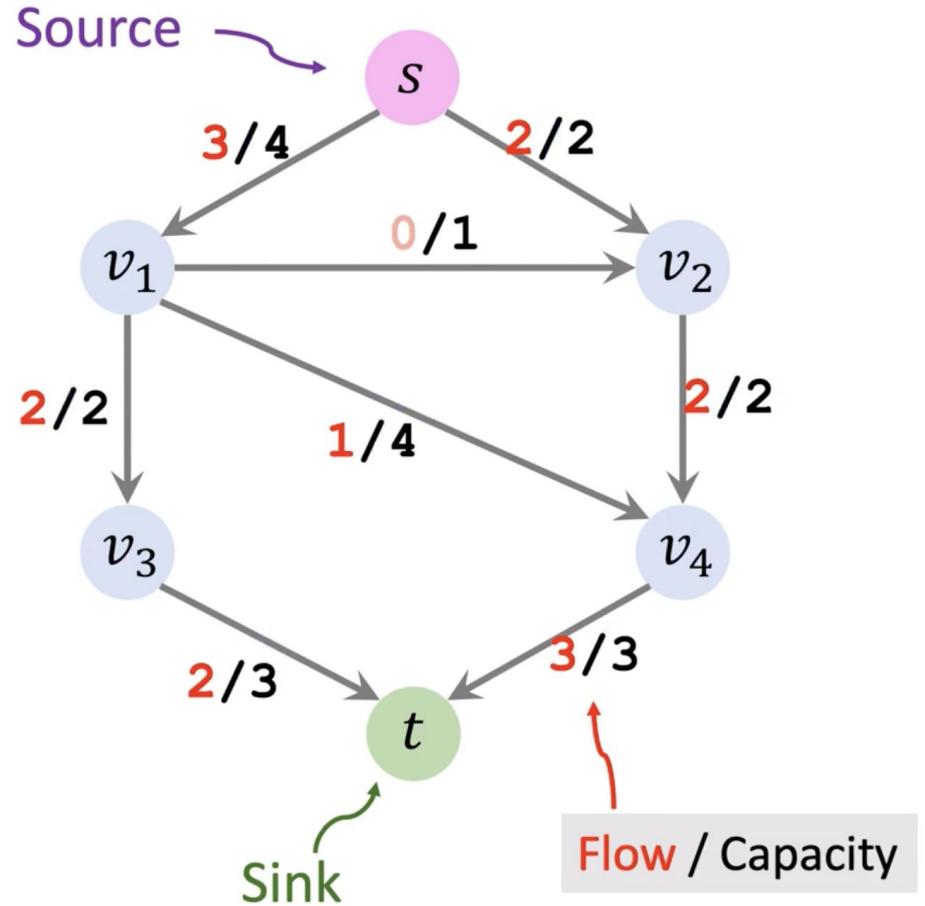
The value $v(f)$ of a flow f is $f^{\text{out}}(s)$.

- That is: it is the amount of material that leaves s .

- Maximum Flow Problem

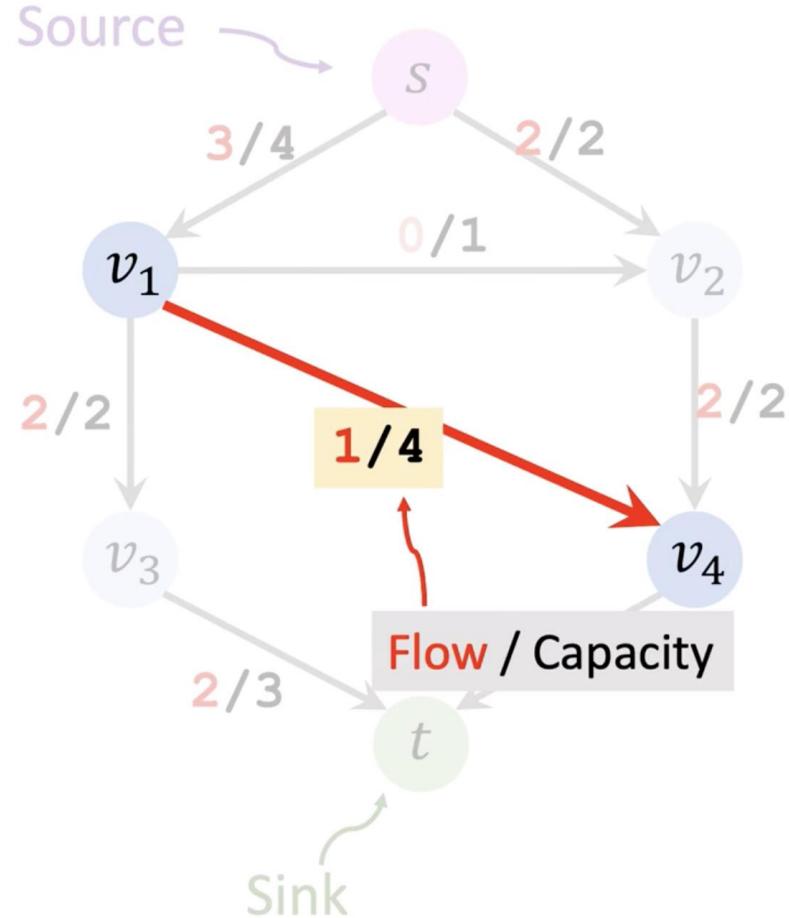
Given a flow network G , find a flow f of maximum possible value.

Example



- Send water from the source s to the sink t .
- The edges are pipes which have certain capacities, e.g., $4 \text{ m}^3/\text{s}$.
- How much water can flow from source s to the sink t at most?

Example



- Capacity of the red pipe is $4 m^3/s$.
- A flow of $1 m^3/s$ goes through the red pipe.
- It has a residual of $3 m^3/s$.

A Naïve Start

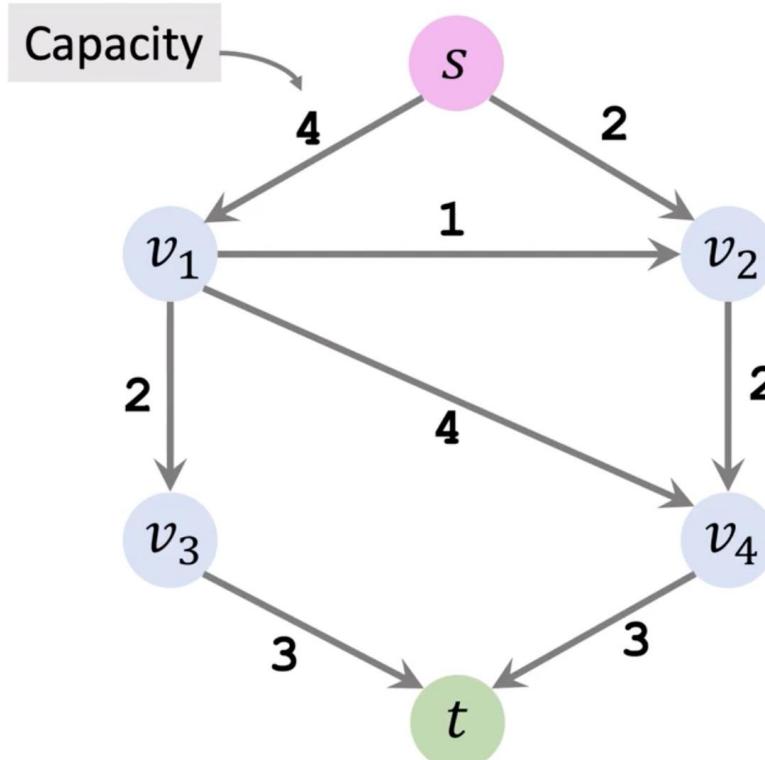
- Suppose we let $f(e) = 0$ for all edges (no flow anywhere).
- Choose some s-t path and “push” flow along it up to the capacities.

Repeat.

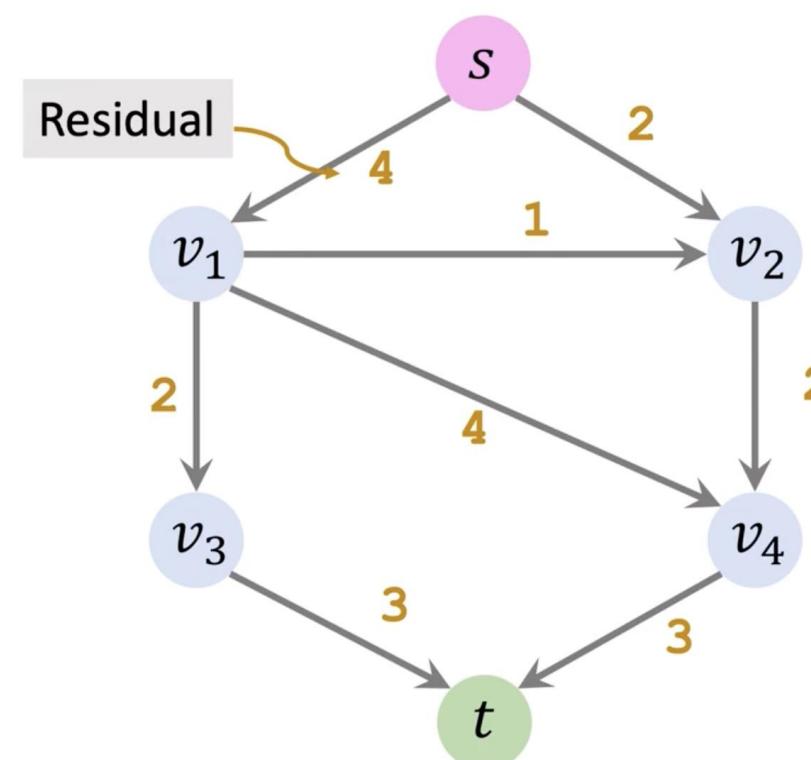
Augmenting Paths

- Let P be an $s-t$ path in the residual graph G_f .
- Let $\text{bottleneck}(P, f)$ be the smallest capacity in G_f on any edge of P .
- If $\text{bottleneck}(P, f) > 0$ then we can increase the flow by sending $\text{bottleneck}(P, f)$ along the path P .

Naïve Example



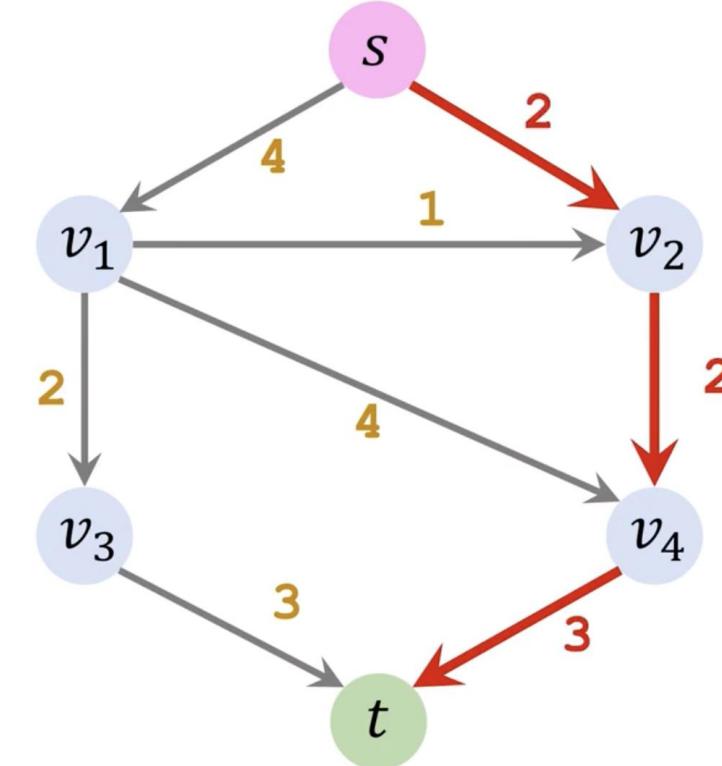
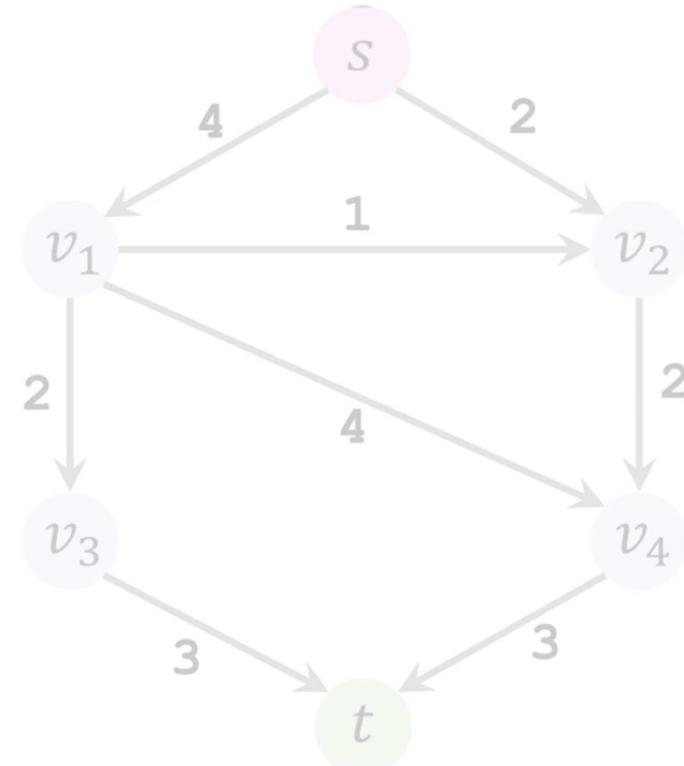
Original Graph



Residual Graph

Naïve Example

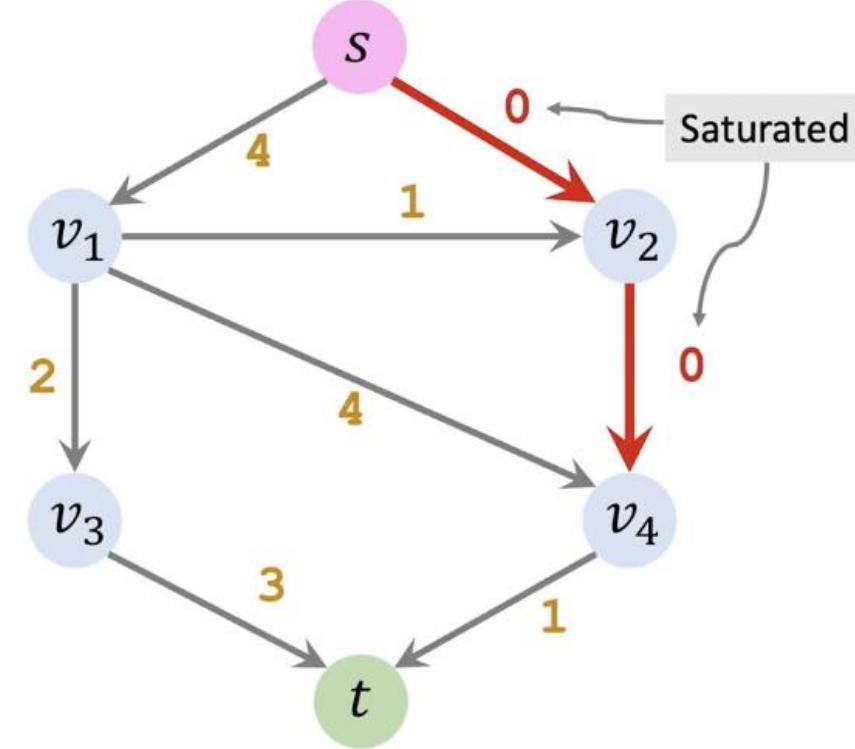
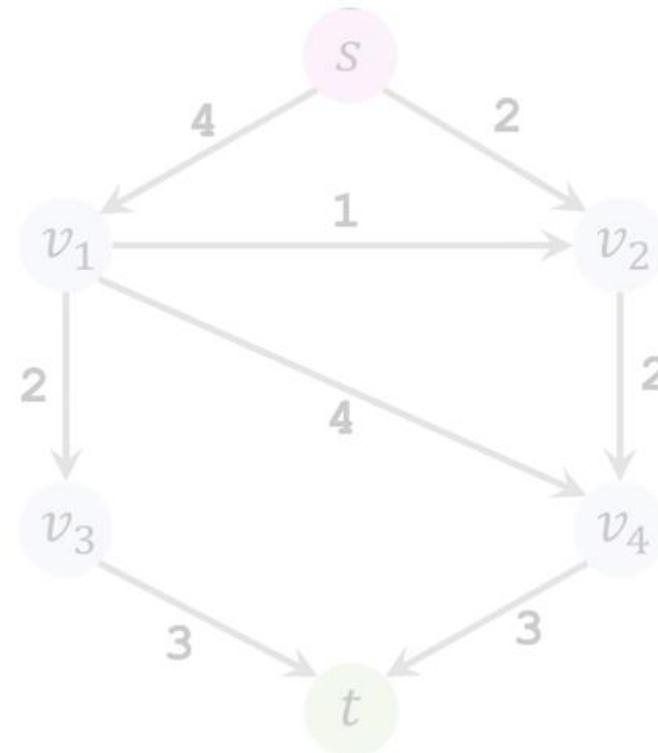
Iteration 1: Find an augmenting path



Found path $s \rightarrow v_2 \rightarrow v_4 \rightarrow t$. (Bottleneck capacity = 2.)

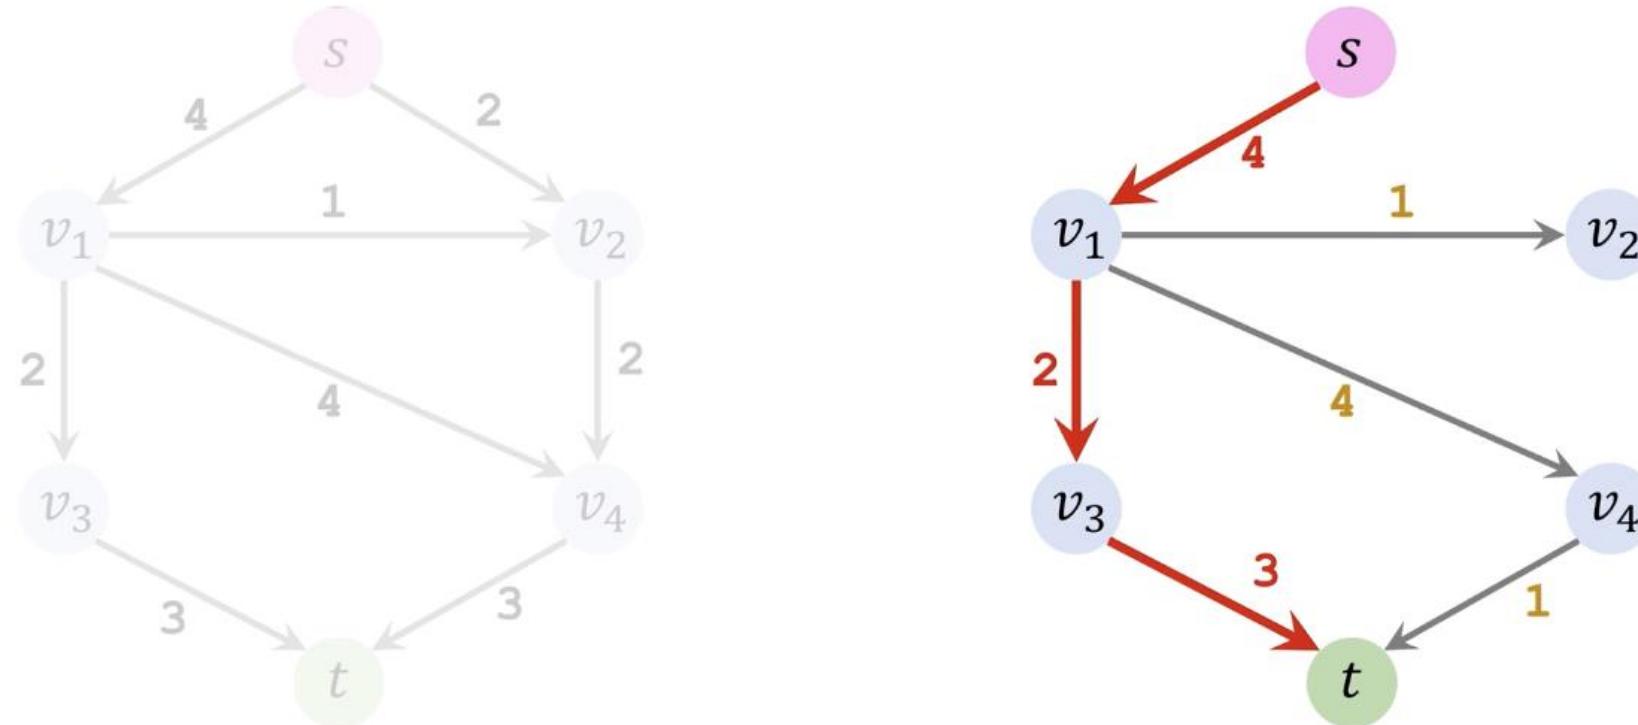
Naïve Example

Iteration 1: Remove saturated edges



Naïve Example

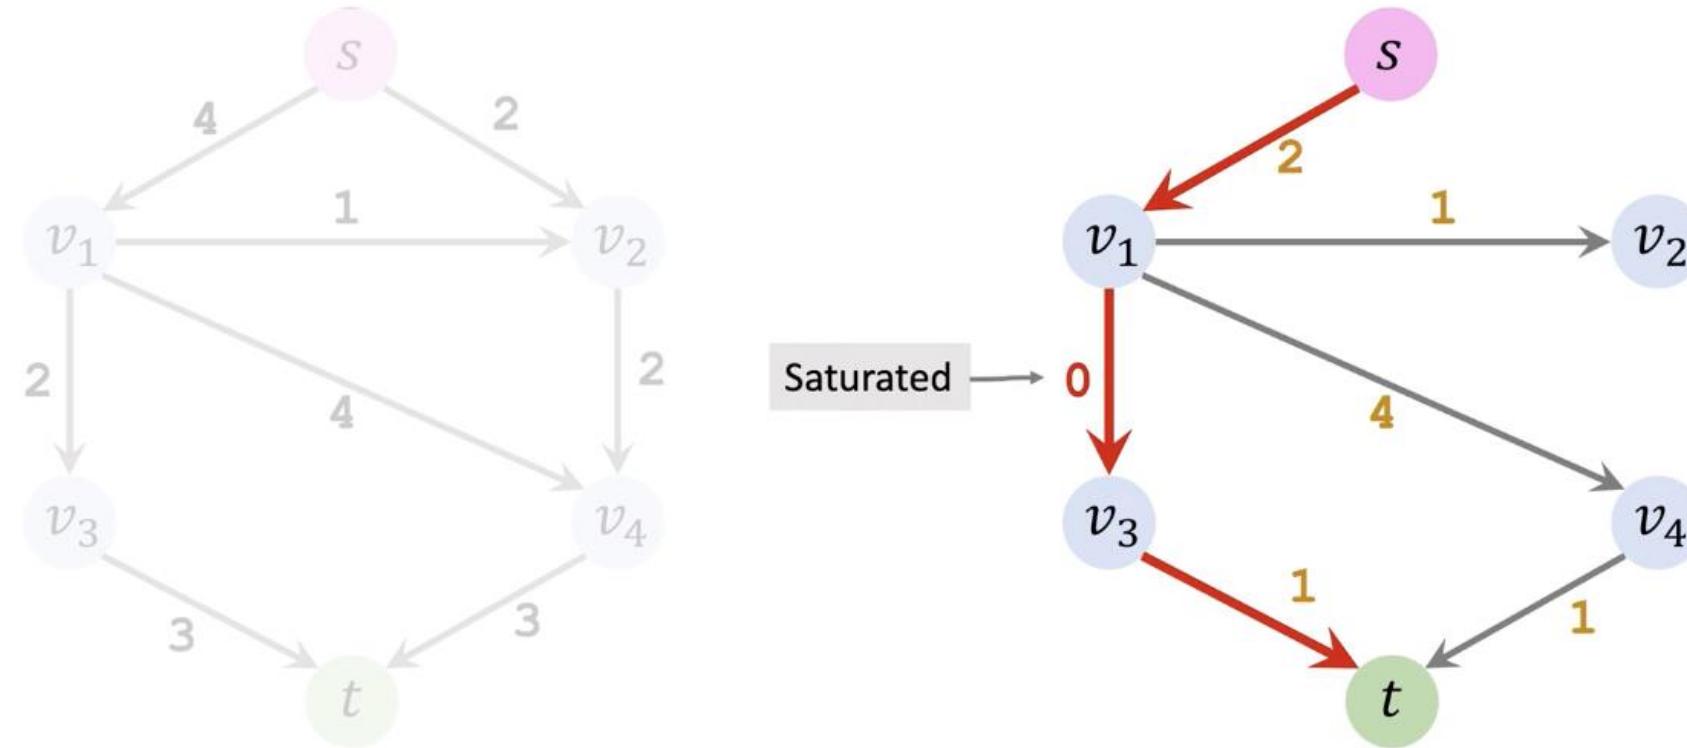
Iteration 2: Find an augmenting path



Found path $s \rightarrow v_1 \rightarrow v_3 \rightarrow t$. (Bottleneck capacity = 2.)

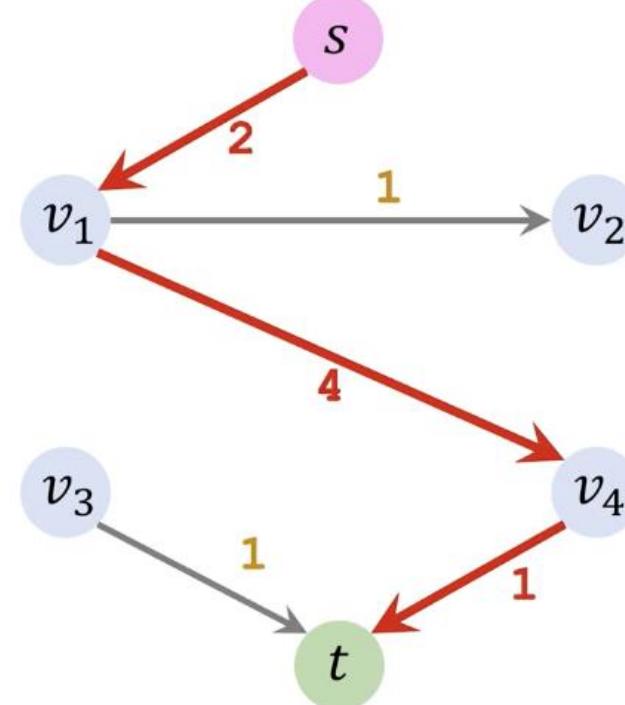
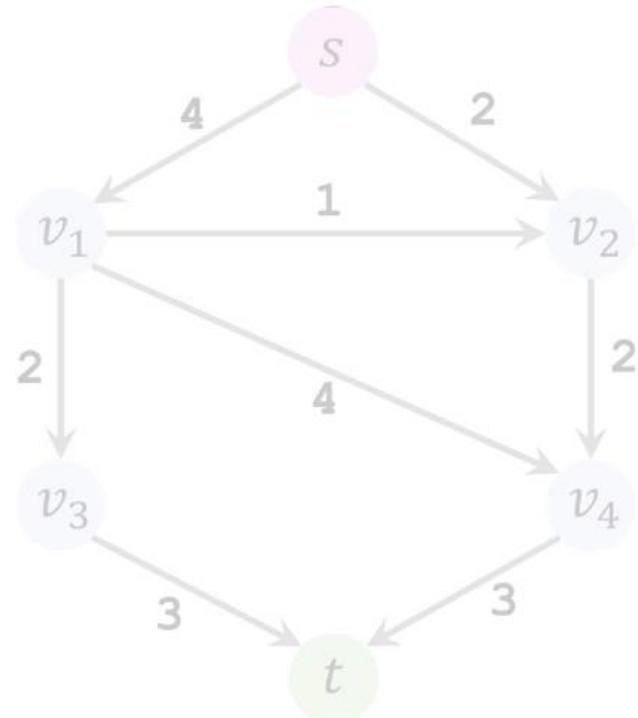
Naïve Example

Iteration 2: Remove saturated edges



Naïve Example

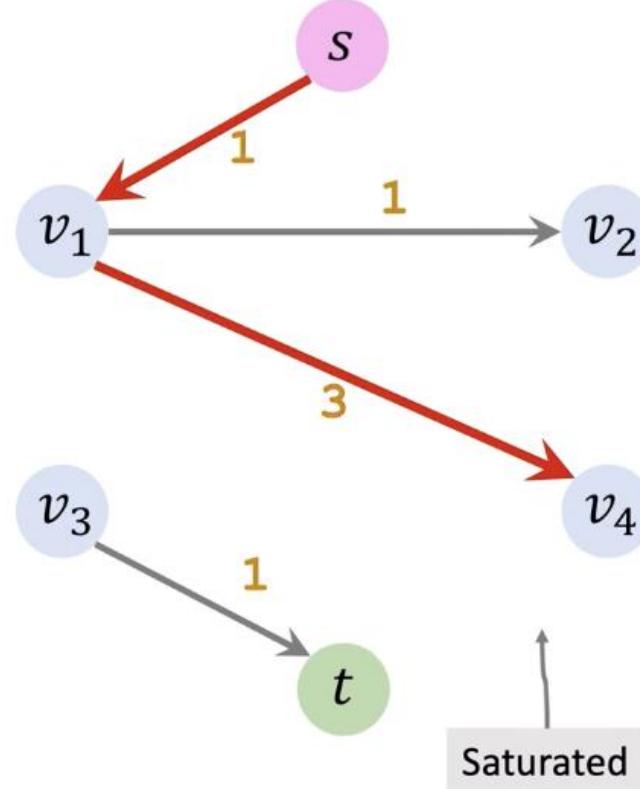
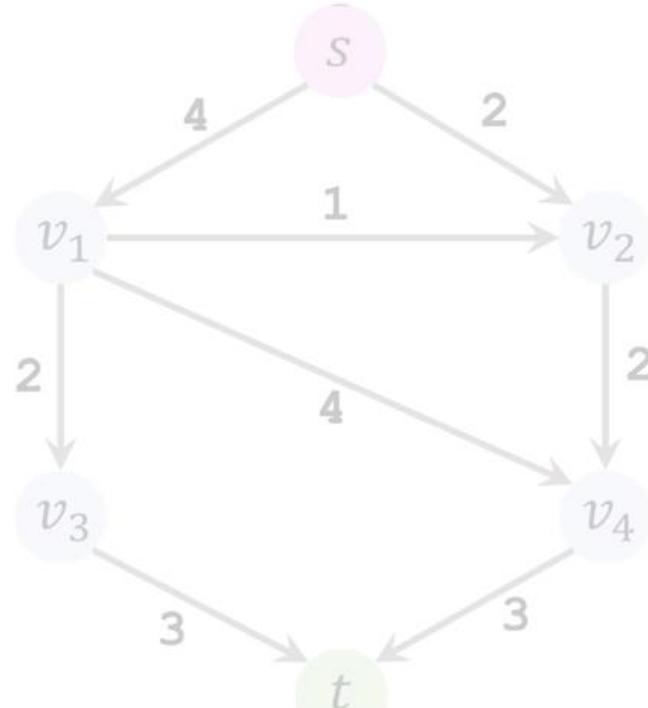
Iteration 3: Find an augmenting path



Found path $s \rightarrow v_1 \rightarrow v_4 \rightarrow t$. (Bottleneck capacity = 1.)

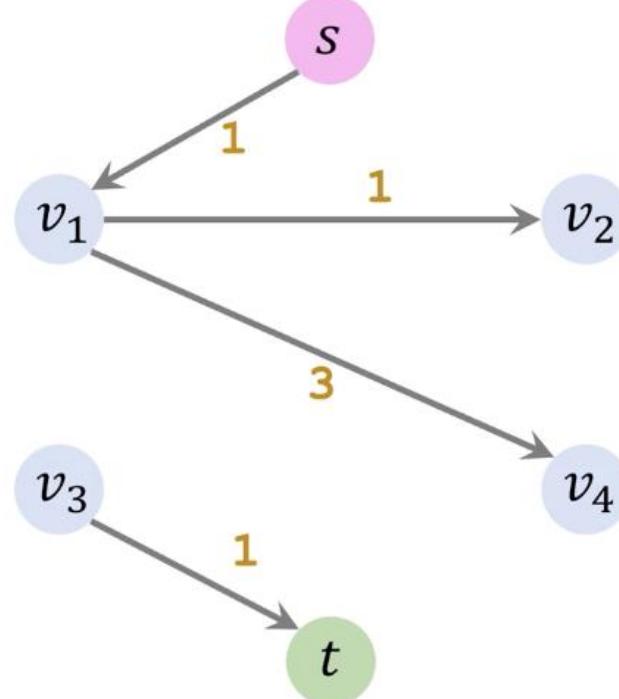
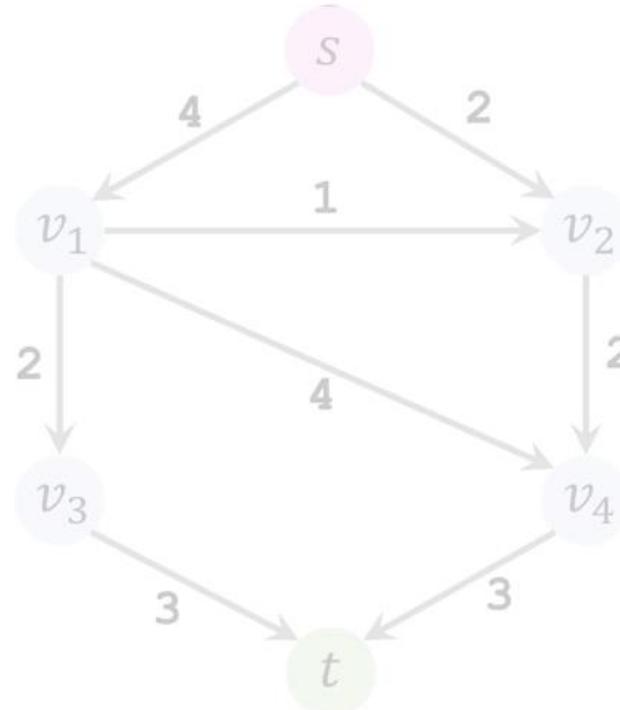
Naïve Example

Iteration 3: Remove saturated edges



Naïve Example

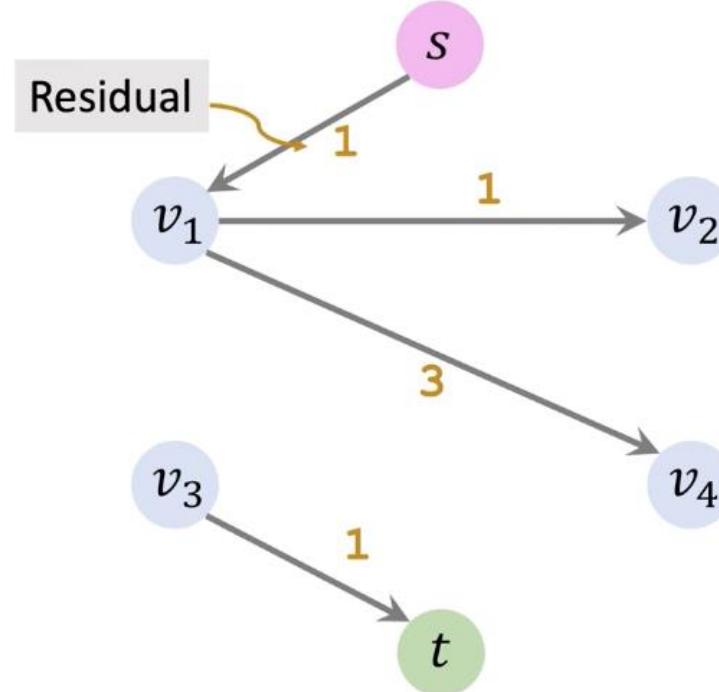
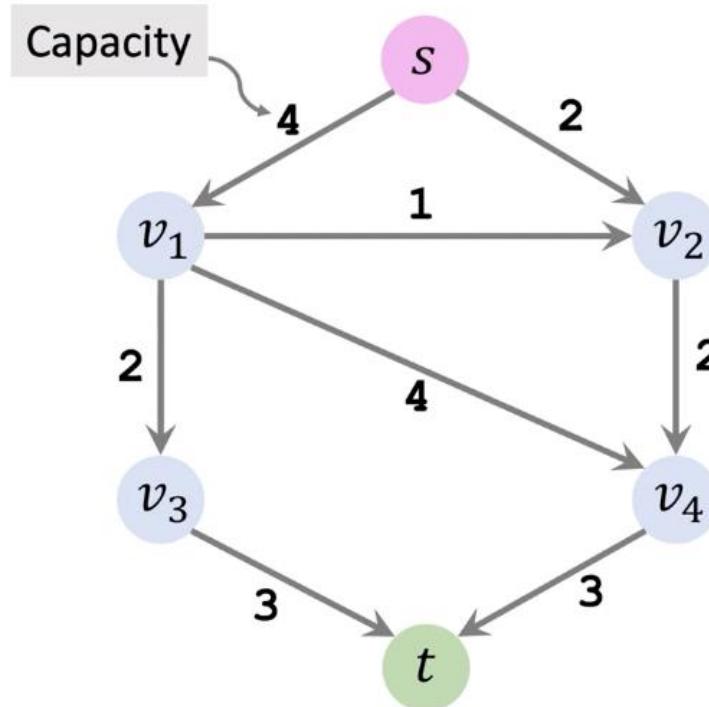
Iteration 4: Find an augmenting path



Cannot find any path from source to sink.

Naïve Example

End of Procedure



$$\text{Flow} = \text{Capacity} - \text{Residual}.$$

Naïve Example

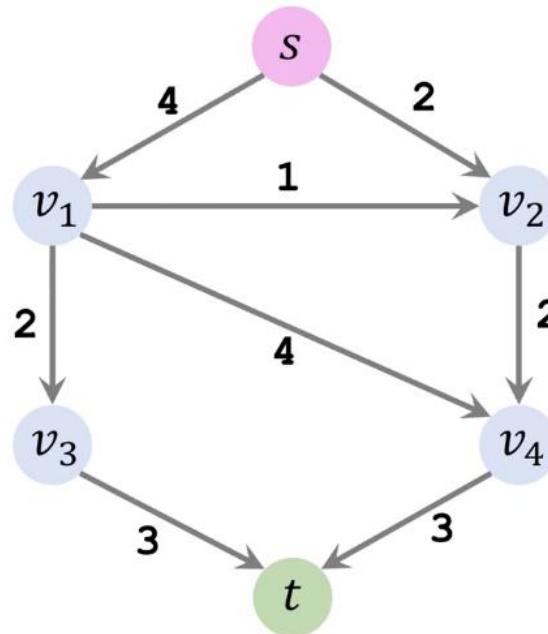
Is the naïve approach always correct?

Problem: Once a bad approach is selected, the naïve approach cannot make any corrections!

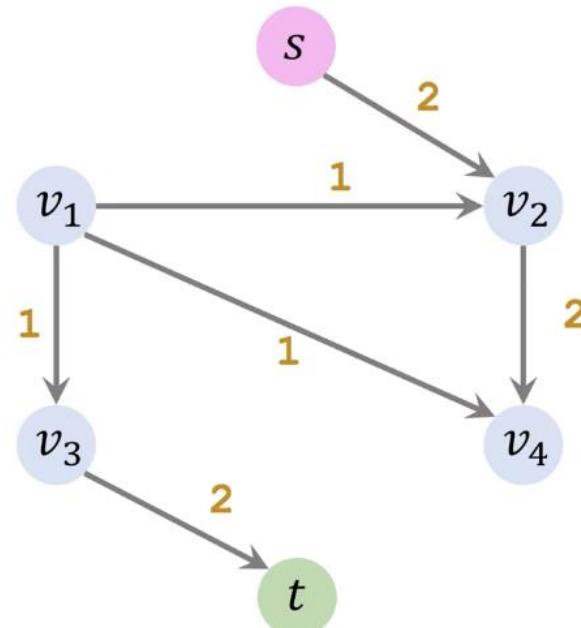
Naïve Example

Consider choosing paths in the order of $\{s, v_1, v_4, t\}$, $\{s, v_1, v_3, t\}$, then we can get the final state as,

End of Procedure



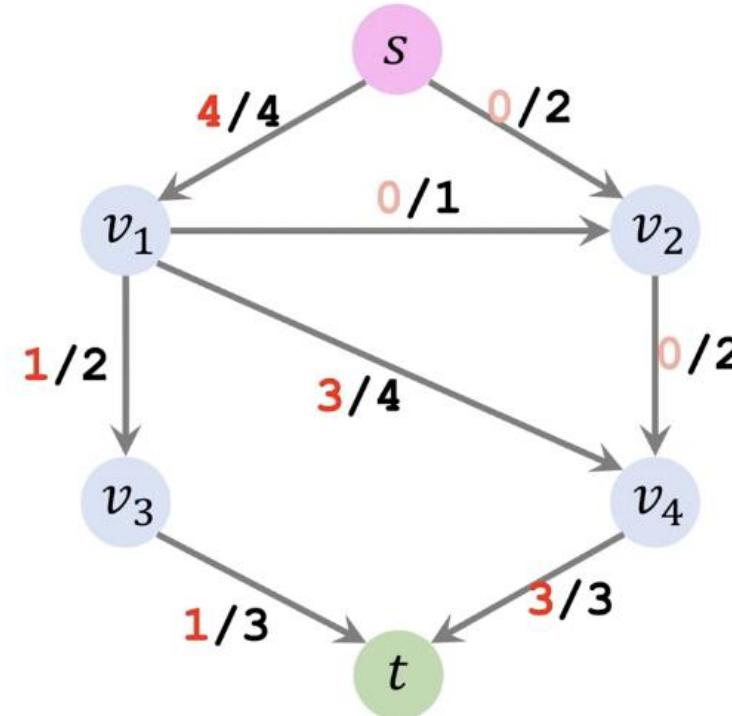
Original Graph



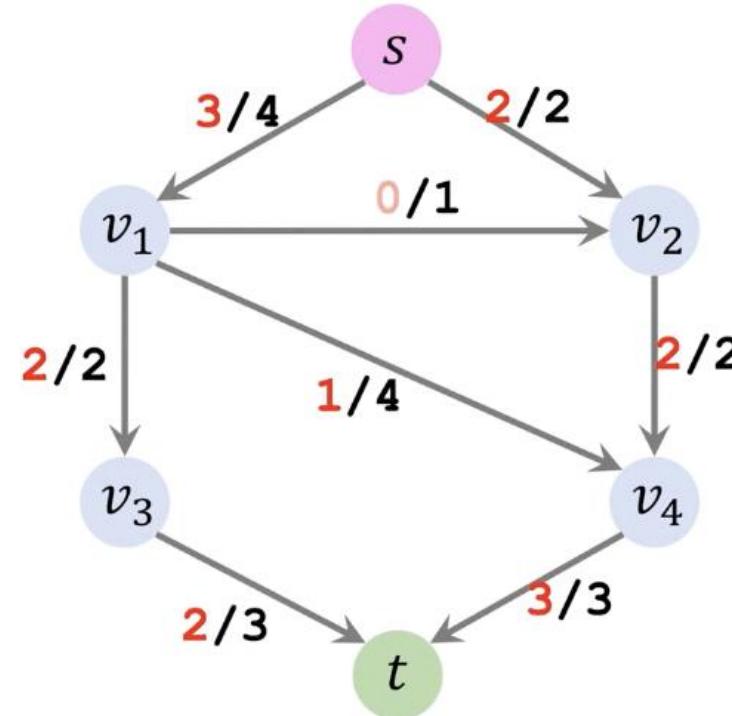
Residual Graph

Naïve Example

The result is not maximum flow!



Flow = 4. (Not maximum!)



Flow = 5. (Maximum!)

Ford-Fulkerson Algorithm

A Greedy Approach

- Suppose we let $f(e) = 0$ for all edges (no flow anywhere).
- Choose some s-t path and “push” flow along it up to the capacities.

Repeat.

- When we get stuck, we can erase some flow along certain edges.

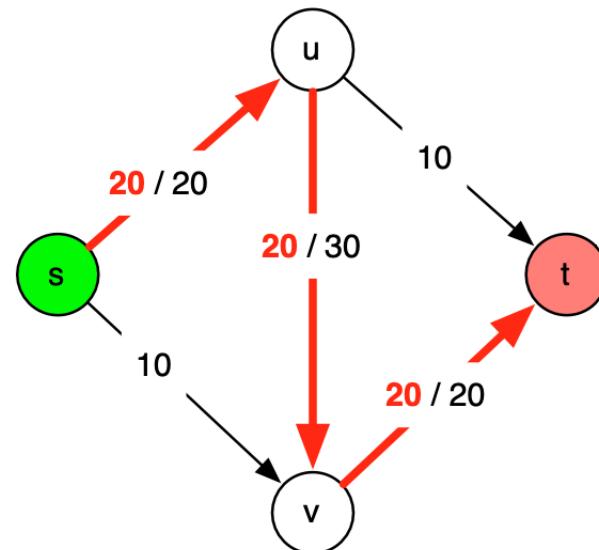
How do we make this more precise?

Residual Graph

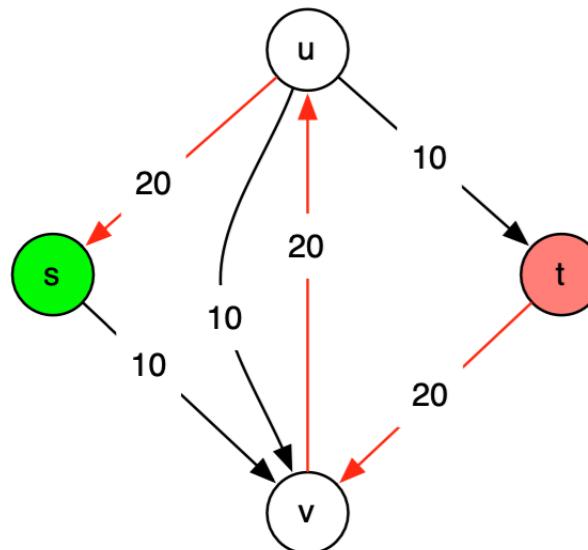
- We define a residual graph G_f . G_f depends on some flow f :
 - G_f contains the same nodes as G .
 - Forward edges: For each edge $e = (u,v)$ of G for which $f(e) < c_e$, include an $e' = (u,v)$ in G_f with capacity $c_e - f(e)$.
 - Backward edges: For each edge $e = (u,v)$ in G with $f(e) > 0$, we include an $e' = (v,u)$ in G_f with capacity $f(e)$.

Residual Graph: Example

- If $f(e) < c_e$, add edge to G_f with capacity $c_e - f(e)$
 - Remaining capacity left
- If $f(u,v) > 0$, add reverse edge (v,u) with capacity $f(e)$
 - Can erase up to $f(e)$ capacity



With Flow f



Residual Graph G_f

Augmenting Paths

- Let P be an $s-t$ path in the residual graph G_f .
- Let $\text{bottleneck}(P, f)$ be the smallest capacity in G_f on any edge of P .
- If $\text{bottleneck}(P, f) > 0$ then we can increase the flow by sending $\text{bottleneck}(P, f)$ along the path P .

Augmenting Paths

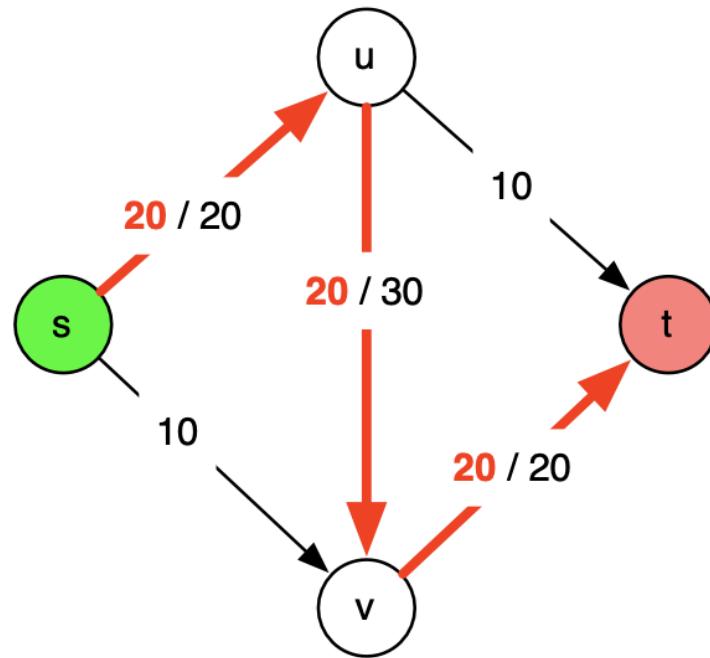
```
augment(f, P):  
    b = bottleneck(P,f)  
    For each edge (u,v) ∈ P:  
        If e = (u,v) is a forward edge:  
            Increase f(e) in G by b //add some flow  
        Else:  
            e' = (v,u)  
            Decrease f(e') in G by b //erase some flow  
        EndIf  
    EndFor  
    Return f
```

Ford-Fulkerson Algorithm

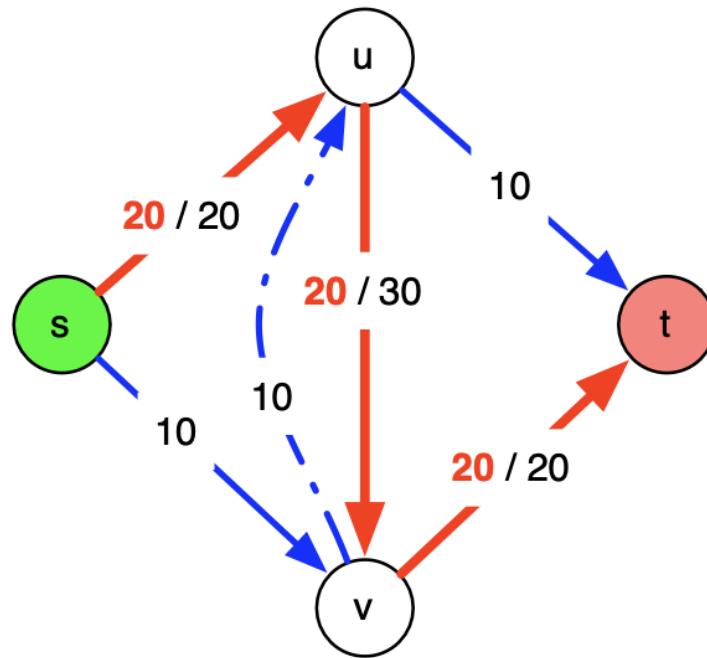
```
MaxFlow(G):
    // initialize:
    Set f[e] = 0 for all e in G

    // while there is an s-t path in Gf:
    While P = FindPath(s,t, Residual(G,f)) != None:
        f = augment(f, P)
        UpdateResidual(G, f)
    EndWhile
    Return f
```

A Glimpse



After 1 path, we've allocated 20 units



Want to send the blue path

After augment, we still have a flow

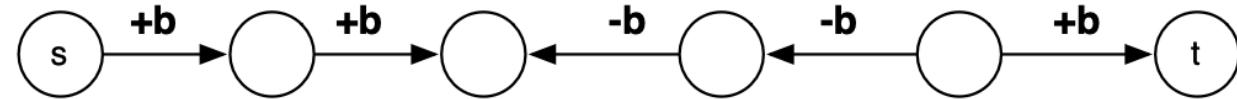
After $f' = \text{augment}(P, f)$, we still have a flow:

Capacity constraints: Let e be an edge on P :

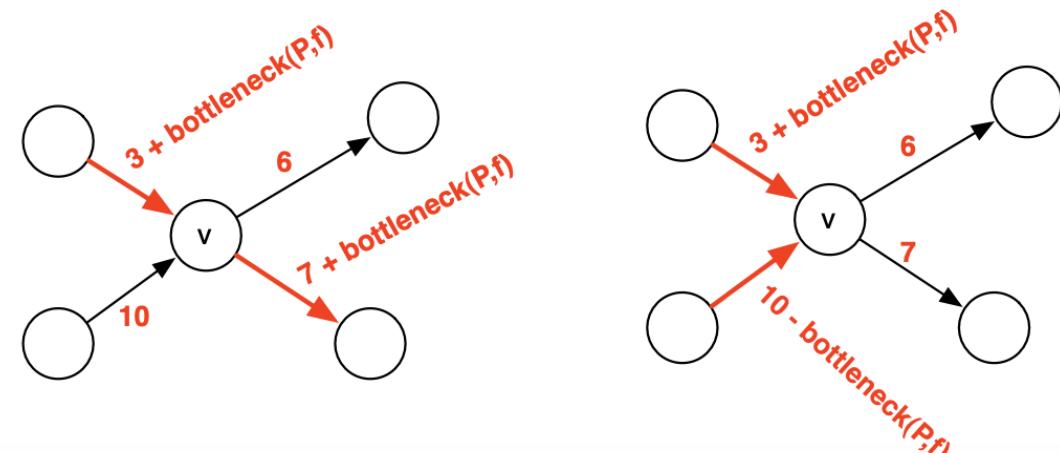
- if e is forward edge, it has capacity $c_e - f(e)$. Therefore,
 - $f'(e) = f(e) + \text{bottleneck}(P, f) \leq f(e) + c_e - f(e) \leq c_e$
- if e is a backward edge, it has capacity $f(e)$. Therefore,
 - $f'(e) = f(e) - \text{bottleneck}(P, f) \geq f(e) - f(e) = 0$

Still have a flow

- Balance constraints: An $s-t$ path in G_f corresponds to some set of edges in G :

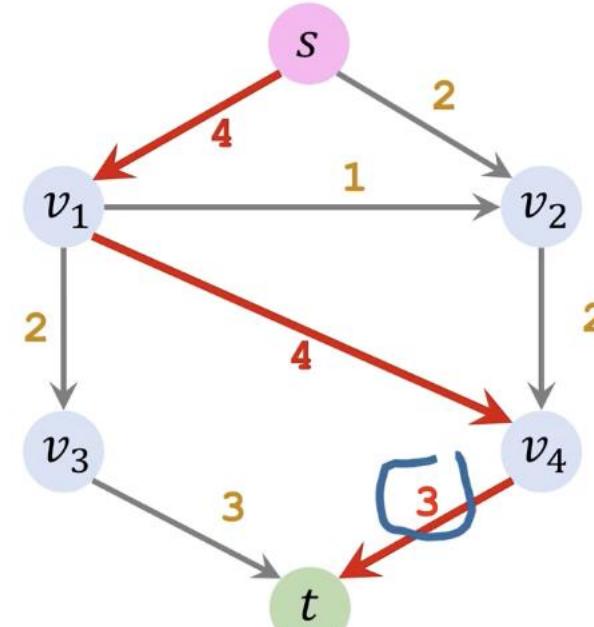
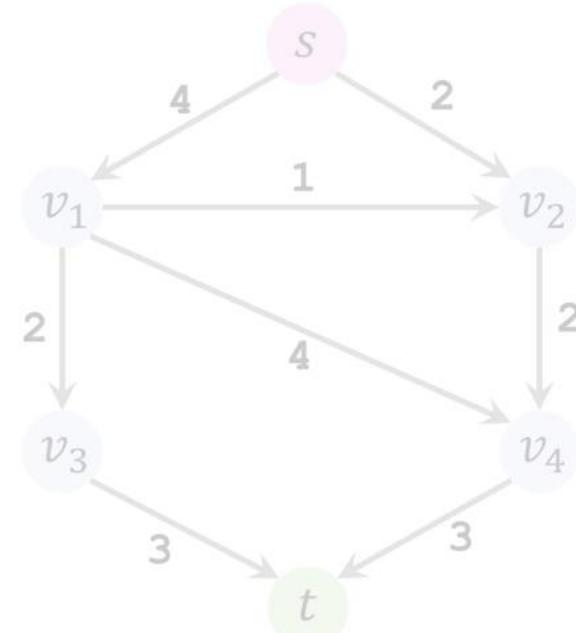


- In other pictures,



Ford-Fulkerson Algorithm: Example

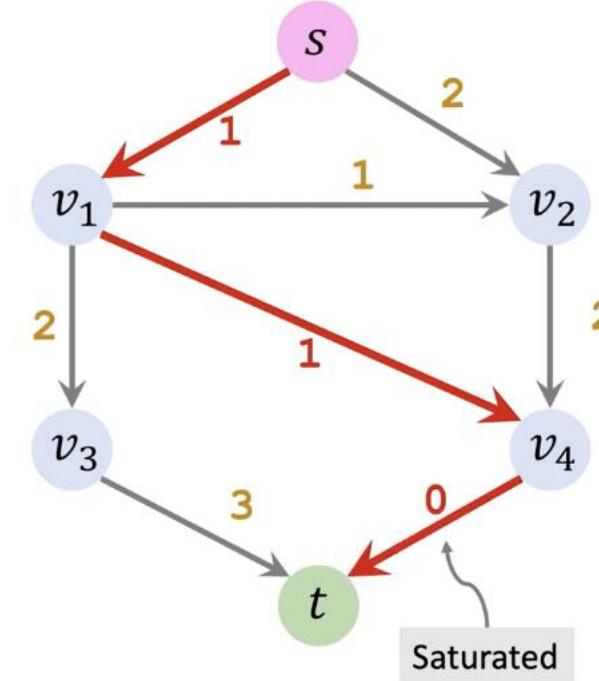
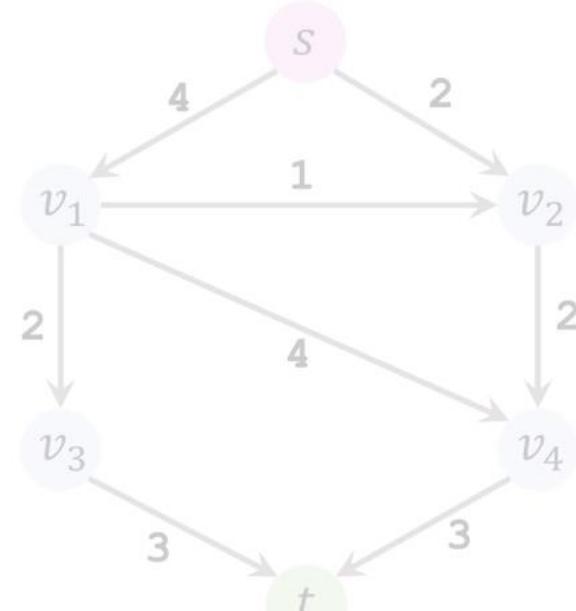
Iteration 1: find an augmenting path



Found path $s \rightarrow v_1 \rightarrow v_4 \rightarrow t$. (Bottleneck capacity = 3.)

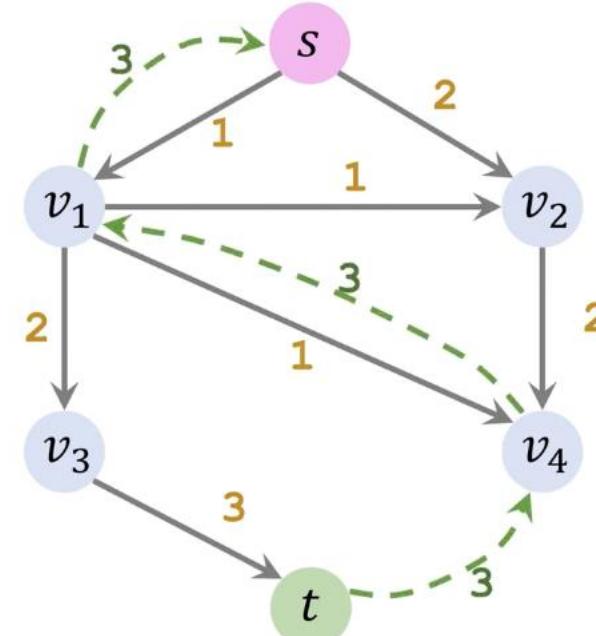
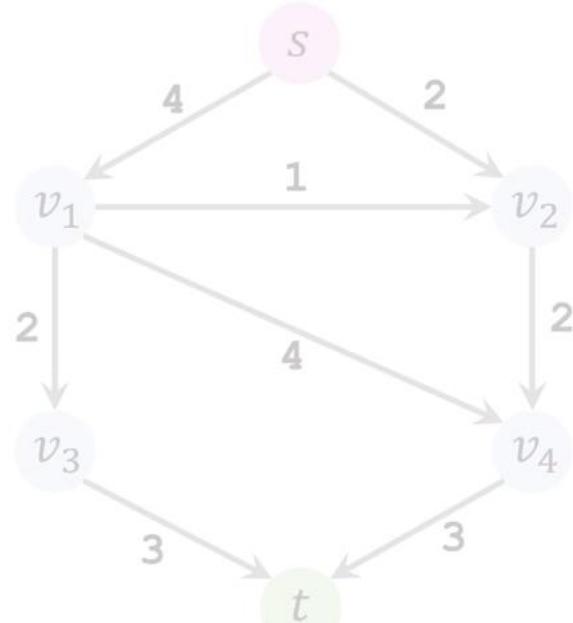
Ford-Fulkerson Algorithm: Example

Iteration 1: update residual graph (remove saturated and add forward edges)



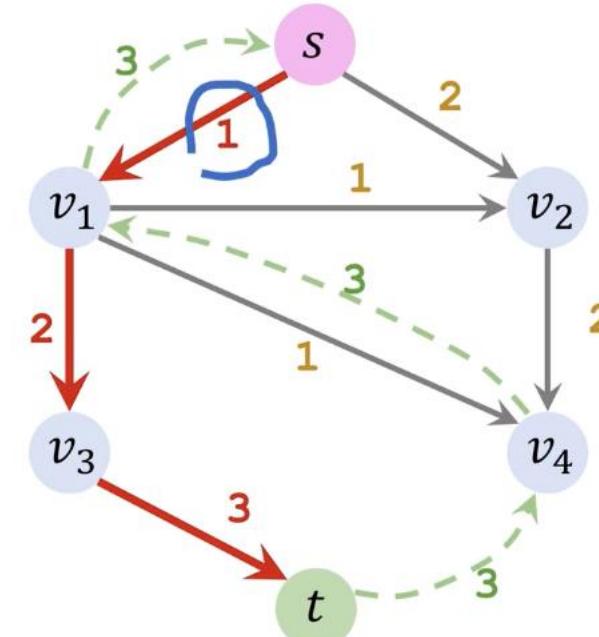
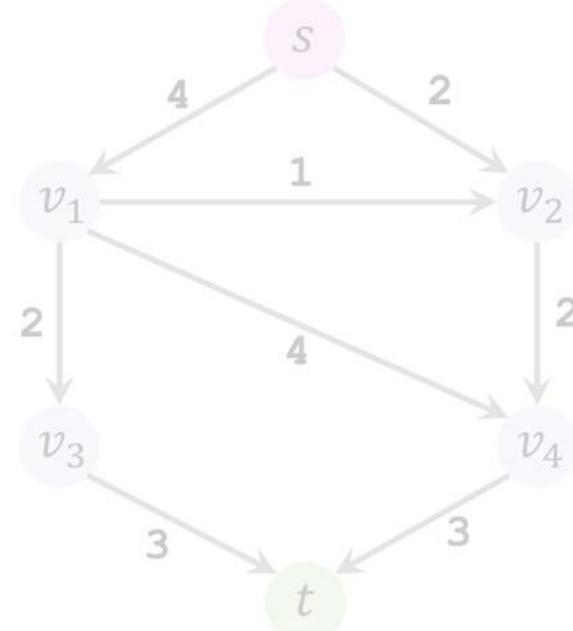
Ford-Fulkerson Algorithm: Example

Iteration 1: update residual graph (add backward edges)



Ford-Fulkerson Algorithm: Example

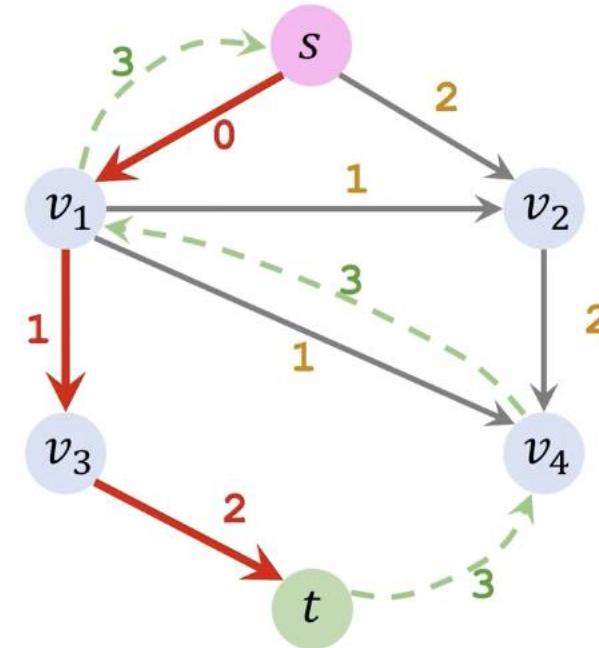
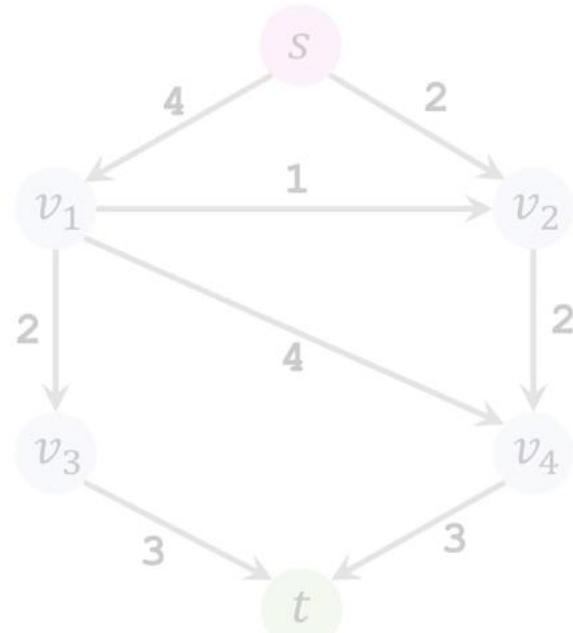
Iteration 2: find an augmenting path



Found path $s \rightarrow v_1 \rightarrow v_3 \rightarrow t$. (Bottleneck capacity = 1.)

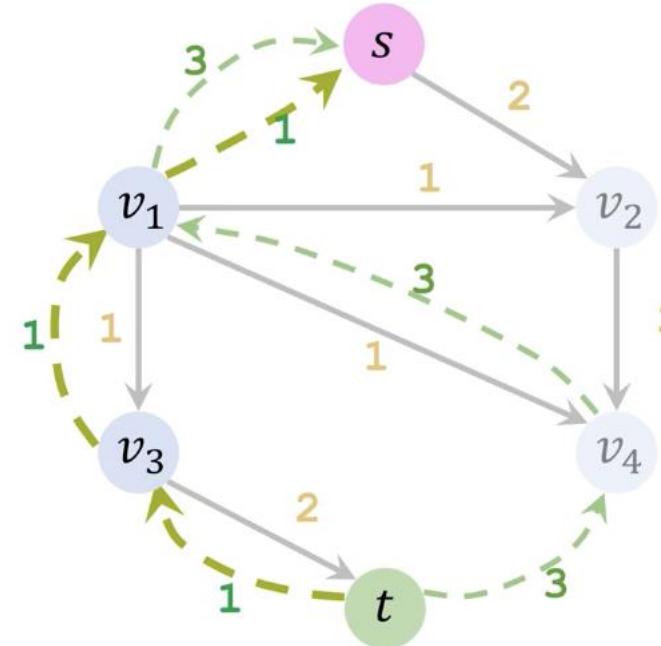
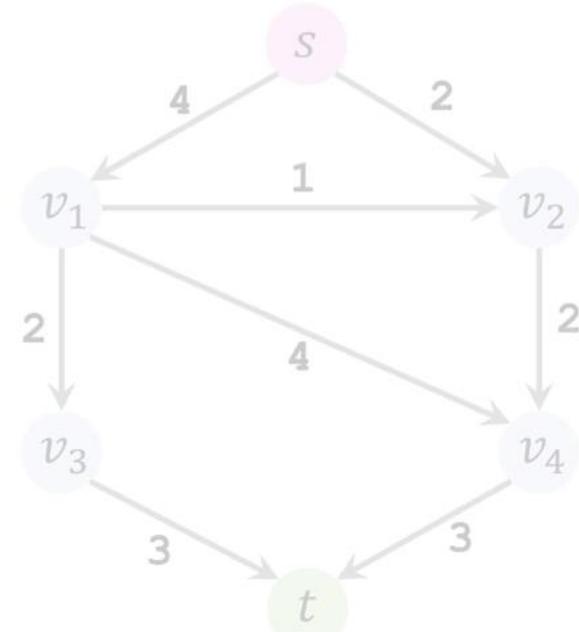
Ford-Fulkerson Algorithm: Example

Iteration 2: update residual graph (remove saturated and add forward edges)



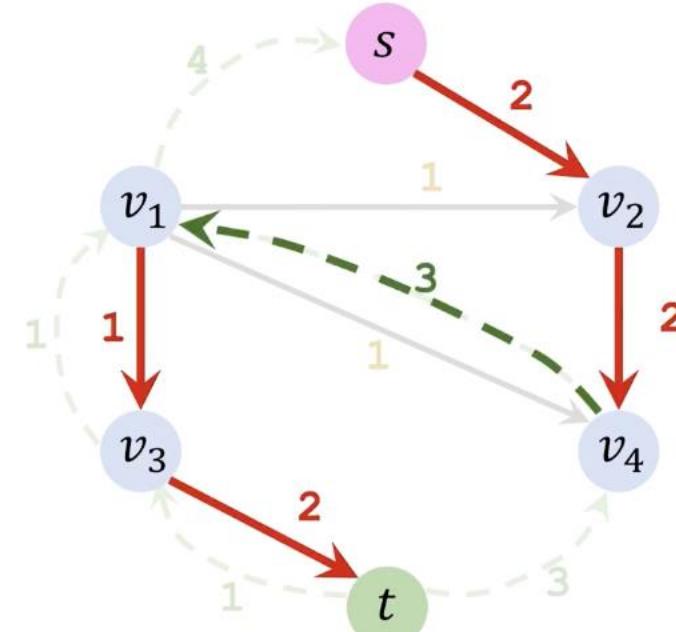
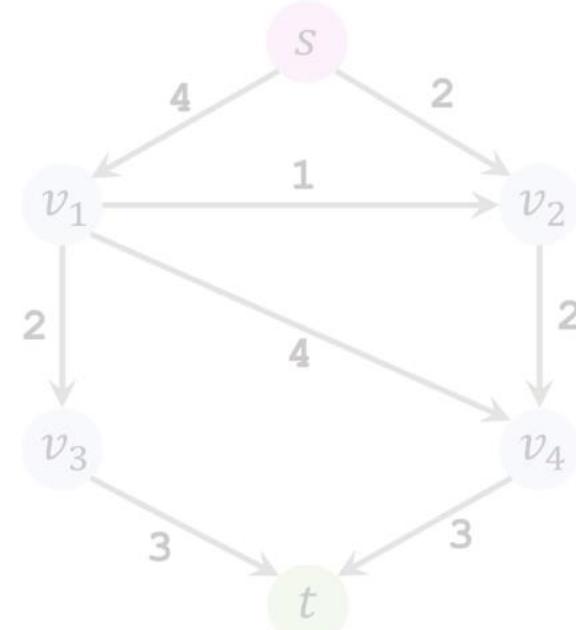
Ford-Fulkerson Algorithm: Example

Iteration 2: update residual graph (add backward edges)



Ford-Fulkerson Algorithm: Example

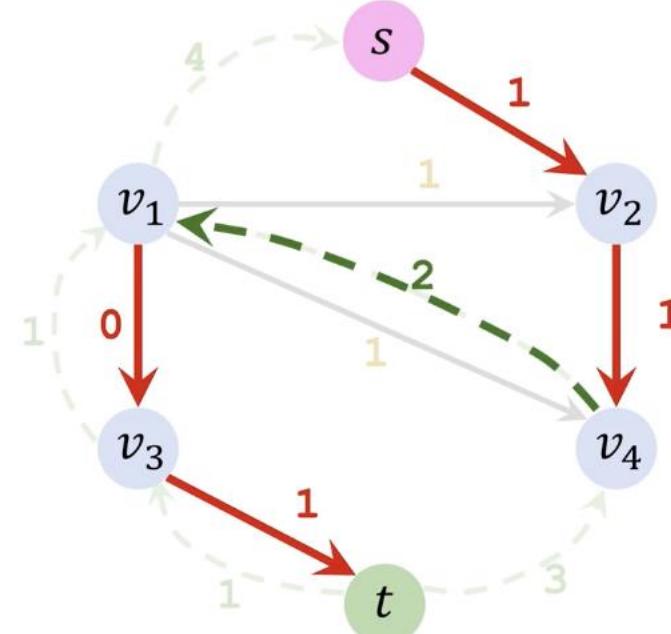
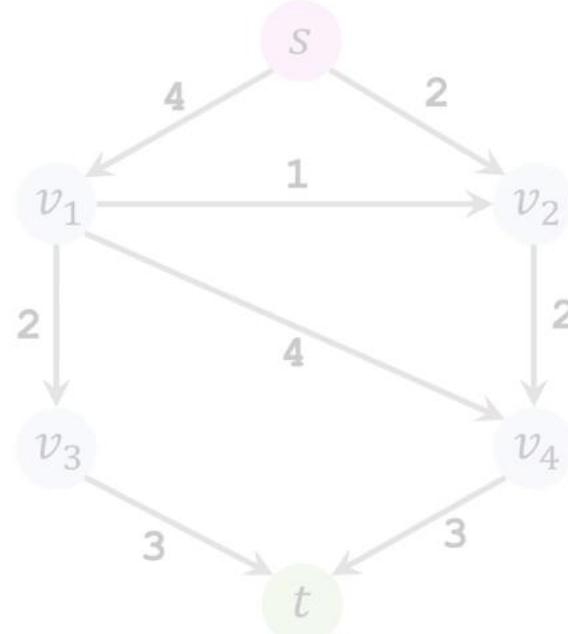
Iteration 3: find an augmenting path



Found path $s \rightarrow v_2 \rightarrow v_4 \rightarrow v_1 \rightarrow v_3 \rightarrow t$.

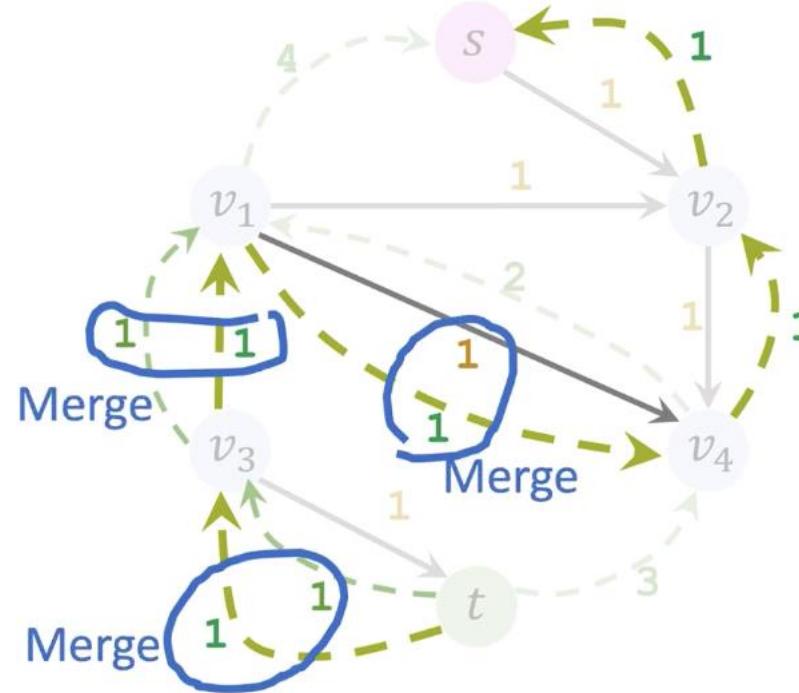
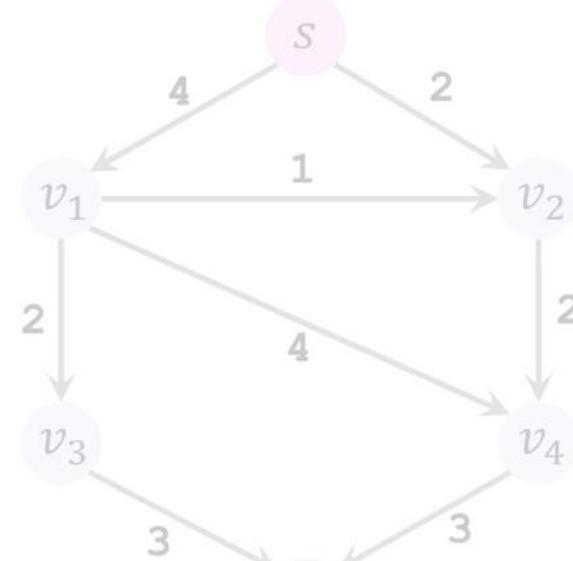
Ford-Fulkerson Algorithm: Example

Iteration 3: update residual graph



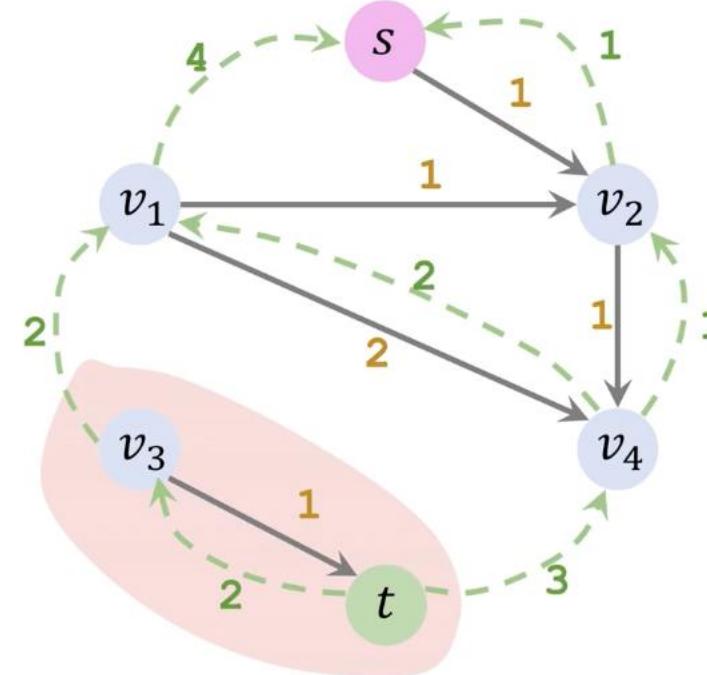
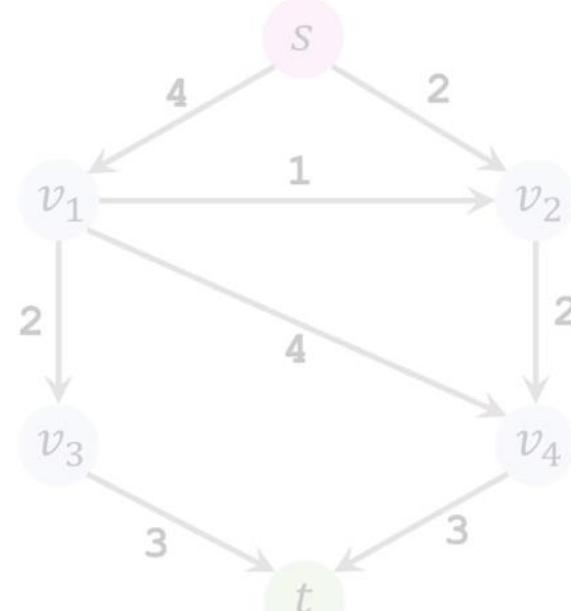
Ford-Fulkerson Algorithm: Example

Iteration 3: update residual graph



Ford-Fulkerson Algorithm: Example

Iteration 4: find an augmenting path

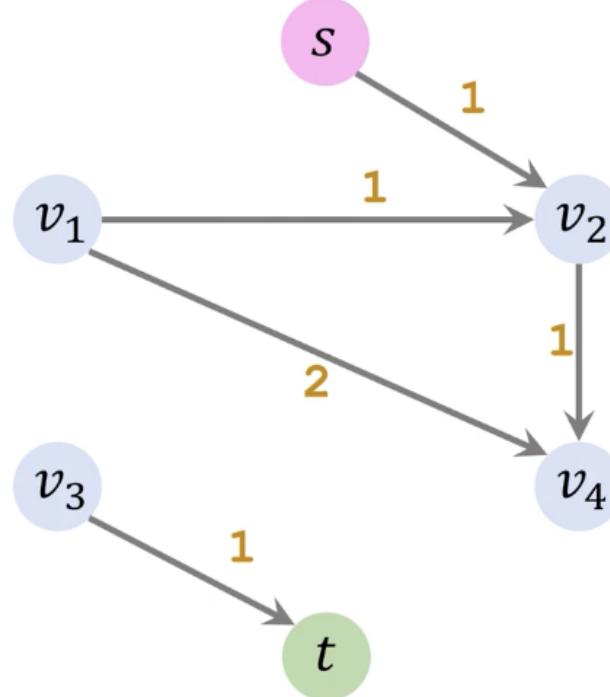
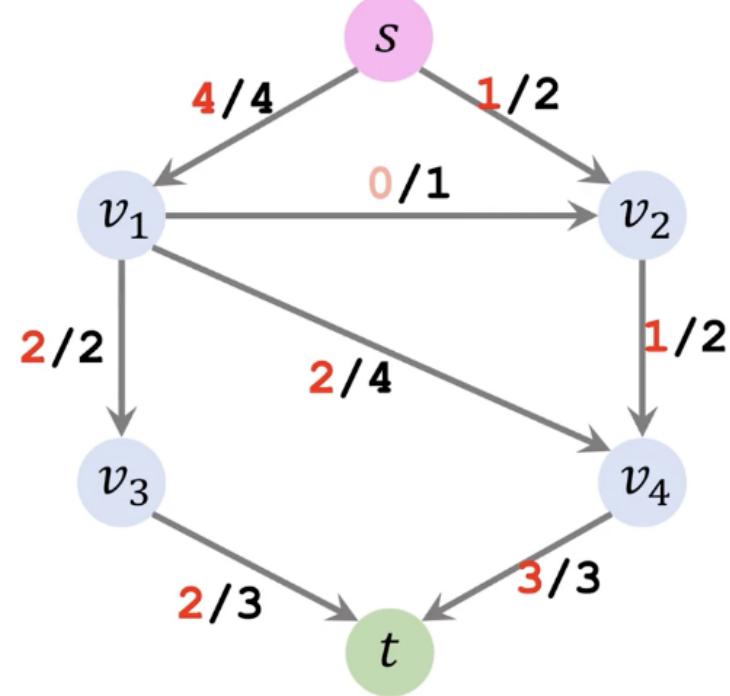


Cannot find any path from source to sink.

Ford-Fulkerson Algorithm: Example

End of Procedure

(backward edges are removed here)



Running Time

- At every step, the flow values $f(e)$ are integers. Start with integers and always add or subtract integers
- At every step we increase the amount of flow $v(f)$ sent by at least 1 unit.
- We can never send more than $C := \sum_e \text{leaving}_s c_e$.
- **Theorem:** The Ford-Fulkerson algorithm terminates in C iterations of the *While* loop.

Time in the While loop

- If G has m edges, G_f has $\leq 2m$ edges.
- Can find an $s-t$ path in G_f in time $O(m+n)$ time with DFS or BFS.
- Since $m \geq n/2$ (every node is adjacent to some edge), $O(m + n) = O(m)$.

Theorem: The Ford-Fulkerson algorithm runs in $O(mC)$ time.

Caveats, etc.

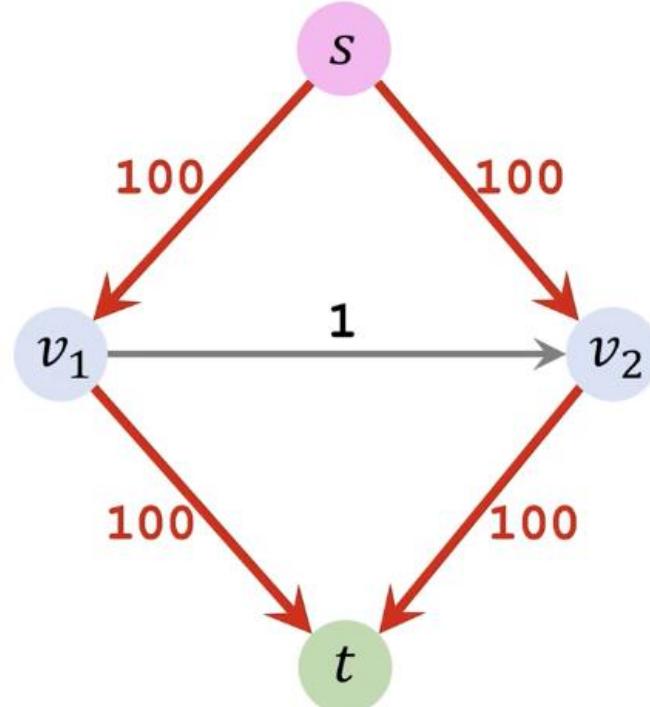
Note this is **pseudo-polynomial** because it depends on the size of the integers in the input.

You can remove this with slightly different algorithms. E.g.:

- $O(nm^2)$: Edmonds-Karp algorithm (use BFS to find the augmenting path)
- $O(m^2 \log C)$ or
- $O(n^2m)$ or $O(n^3)$

Ford-Fulkerson Algorithm: Example

A bad case for Ford-Fulkerson algorithm



- The amount of the max flow is 200.
- Ford-Fulkerson algorithm may take 200 iterations to find the max flow.

Edmonds-Karp Algorithm

Overview

- Edmonds-Karp algorithm [2] is almost the same as as Ford-Fulkerson algorithm [1].
- Edmonds-Karp algorithm uses the shortest path from source to sink.
(Apply weight 1 to all the edges of the residual graph.)
- Edmonds-Karp algorithm is a special case of Ford-Fulkerson algorithm.

Reference

1. L. R. Ford and D. R. Fulkerson. [Maximal flow through a network](#). *Canadian Journal of Mathematics*, 8: 399–404, 1956.
2. J. Edmonds and R. M. Karp. [Theoretical improvements in algorithmic efficiency for network flow problems](#). *Journal of the ACM*. 19 (2): 248–264, 1972.

Ford-Fulkerson Algorithm

```
MaxFlow(G):
    // initialize:
    Set f[e] = 0 for all e in G

    // while there is an s-t path in Gf:
    While P = FindPath(s,t, Residual(G,f)) != None:
        f = augment(f, P)
        UpdateResidual(G, f)
    EndWhile
    Return f
```

Time complexity: $O(Cm)$. (C is the max flow; m is #edges.)

Edmonds-Karp Algorithm

- Follow the same procedure as Ford-Fulkerson, except that we find **the shortest augmenting path** each time (on the residual graph).
- When finding paths, regard the residual graph as unweighted.
- Time complexity: $O(nm^2)$. (m is #edges; n is #vertices.)

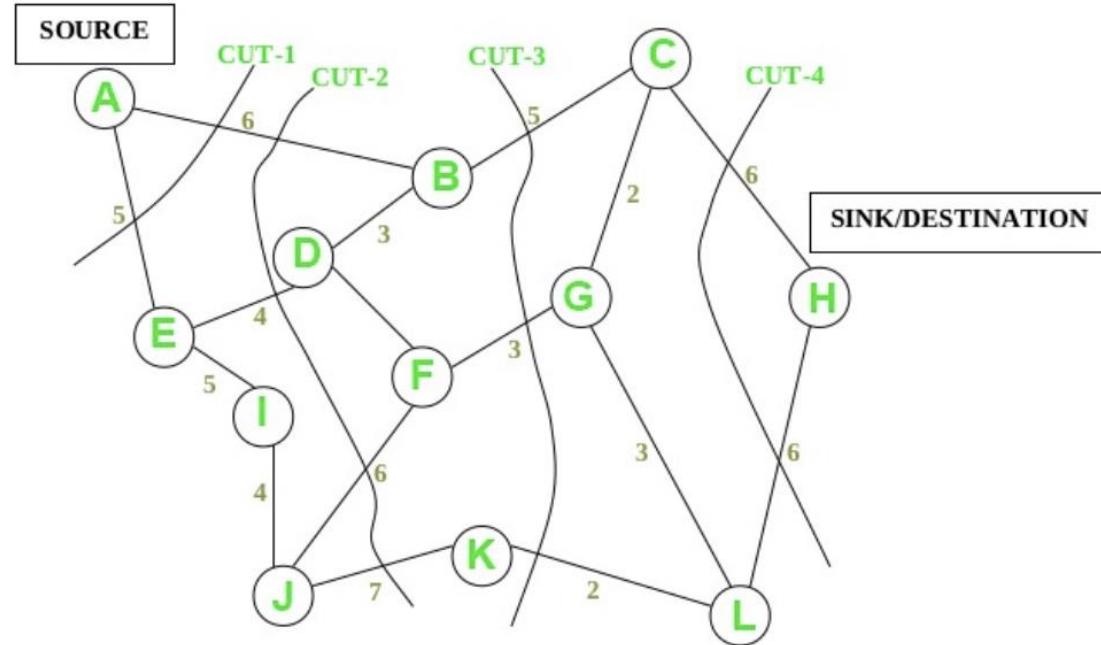
Time Complexity Analysis

- m : number of edges.
- n : number of vertices.
- Each iteration has $O(m)$ time complexity.
- The number of iterations is at most $m \cdot n$. 
- The worst-case time complexity is $O(nm^2)$. Prove by yourself with ref2 if you are interested

Proof of Correctness

How do we know the flow is maximum?

Cuts in graphs



- What is the capacity of a cut? What about if edges are directed?

CUT-1 : AB, AE

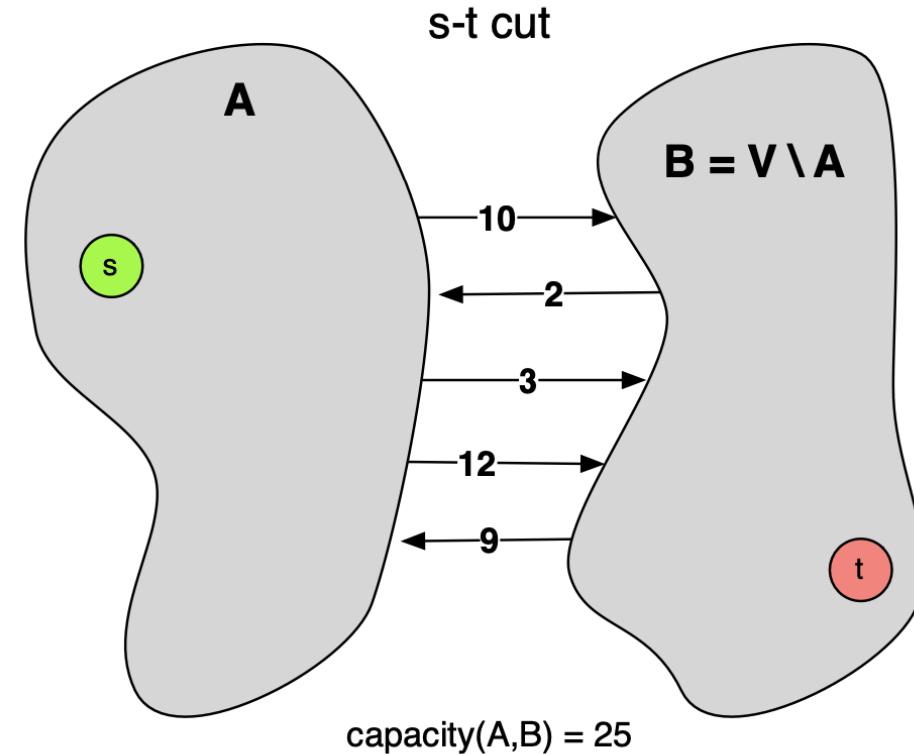
CUT-2 : AB, ED, JF, JK

CUT-3 : BC, FG, KL

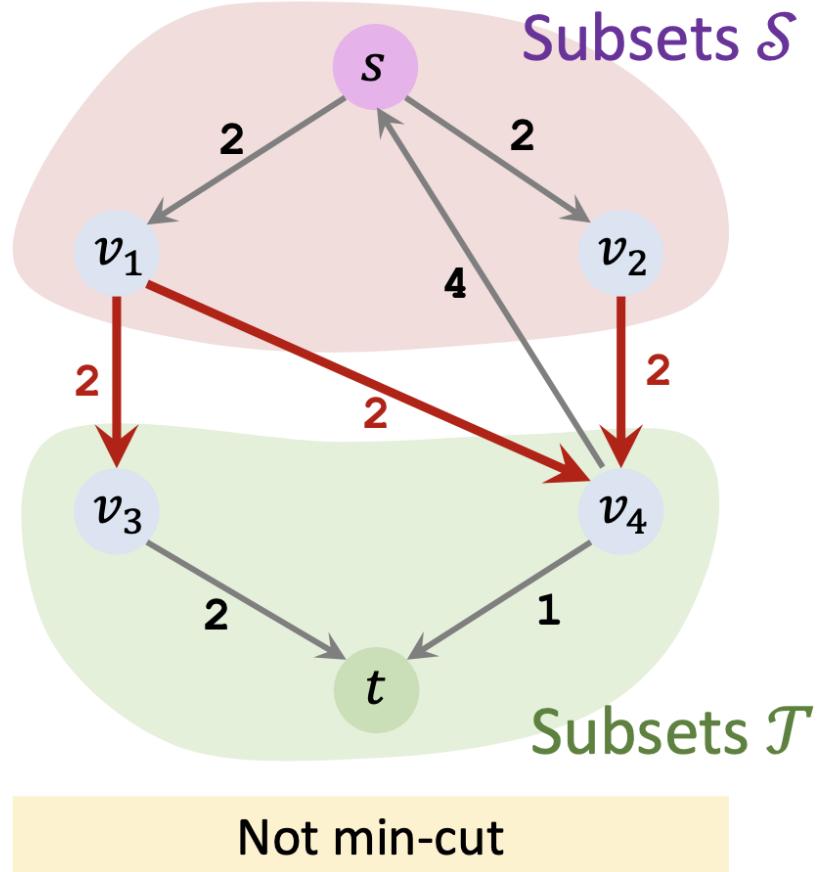
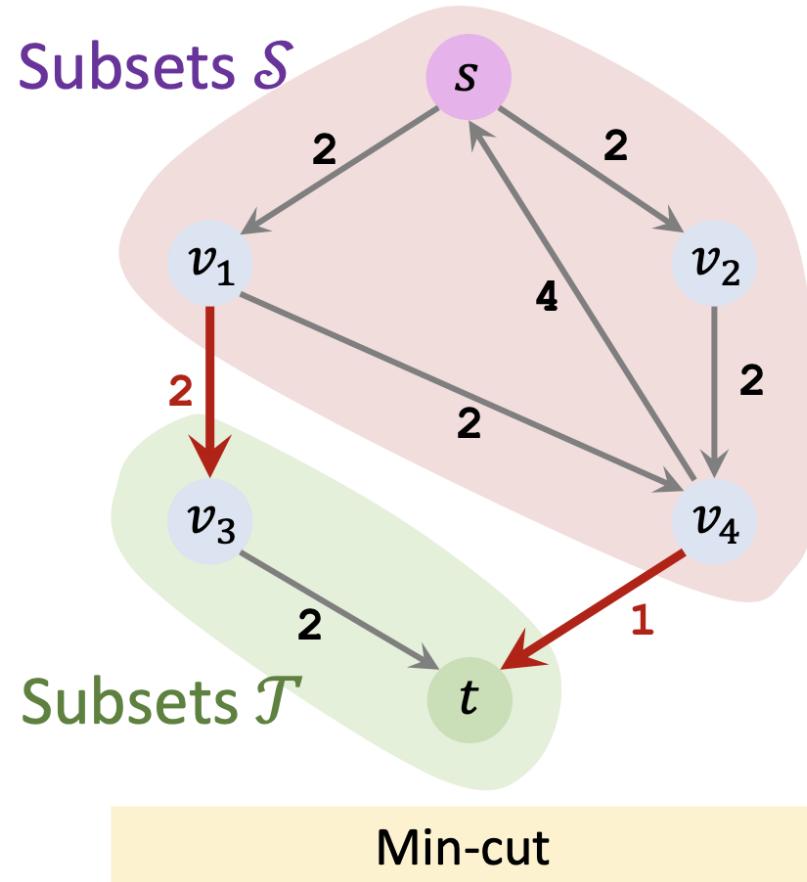
CUT-4 : CH, HL

Cuts and Cut Capacity

- Definition: The capacity (A, B) of an $s-t$ cut (A, B) is the sum of the capacities of edges leaving A .
- Definition: Min-Cut is the $s-t$ cut that minimizes capacity (A, B) .

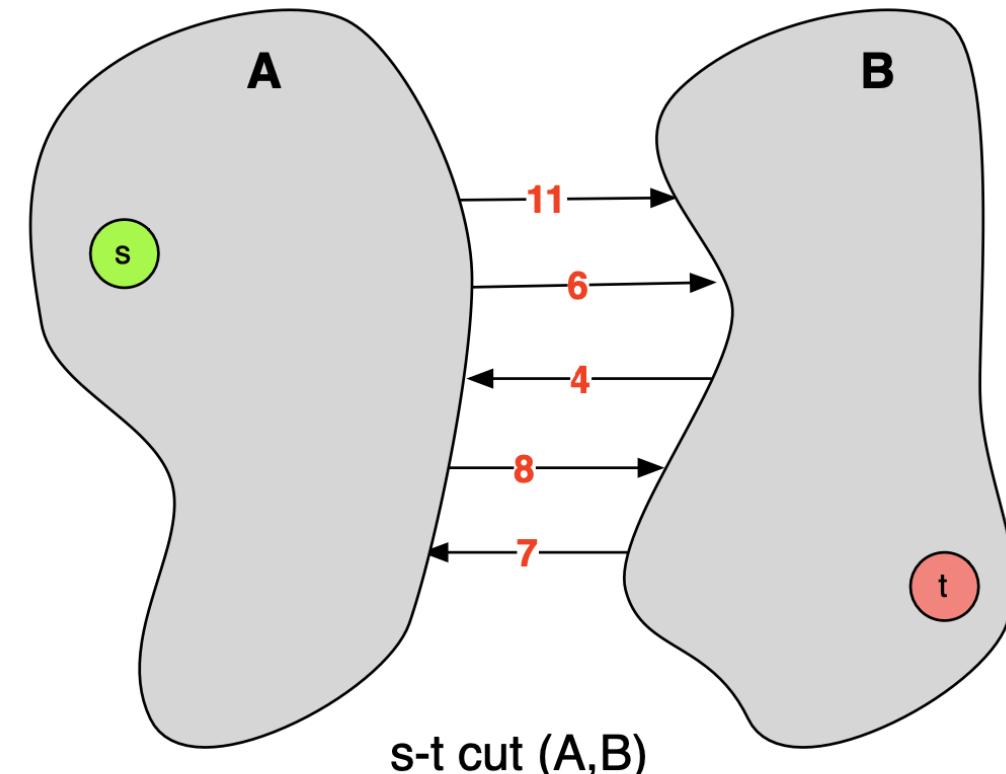


Example



A Cut Theorem

Theorem: Let f be an s - t flow and (A, B) be an s - t cut. Then $v(f) = f^{\text{out}}(A) - f^{\text{in}}(A)$.



Another cut theorem

- Theorem: Let f be any s-t flow and (A, B) be any s-t cut. Then $v(f) \leq \text{capacity}(A, B)$.

Proof.

$$\begin{aligned}
 v(f) &= f^{\text{out}}(A) - f^{\text{in}}(A) && \text{prev. thm} \\
 &\leq f^{\text{out}}(A) && f^{\text{in}}(A) \text{ is } \geq 0 \\
 &= \sum_{e \text{ leaving } A} f(e) && \text{by definition} \\
 &\leq \sum_{e \text{ leaving } A} c_e && \text{by capacity constraints} \\
 &= \text{capacity}(A, B) && \text{by definition}
 \end{aligned}$$

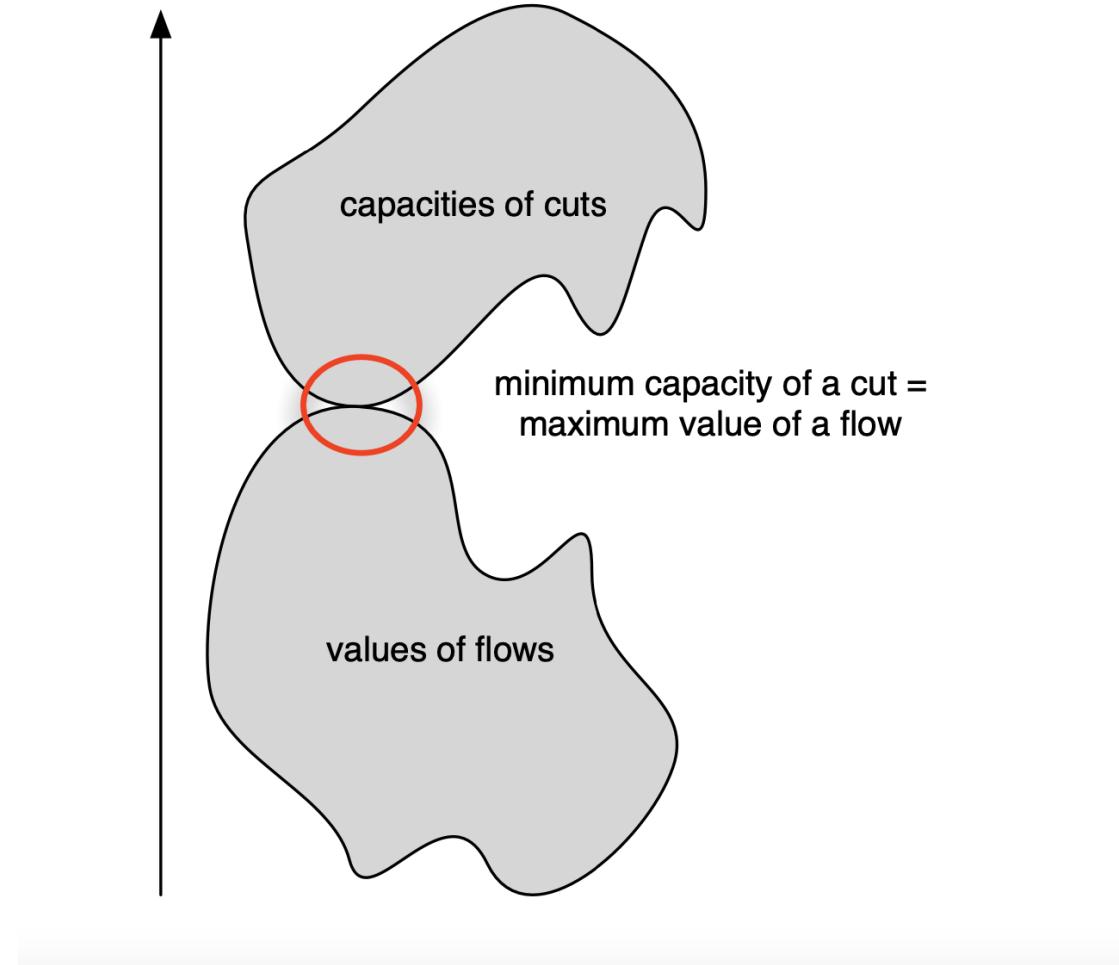


Cuts Constrain Flows

Theorem: Let f be any s - t flow and (A, B) be any s - t cut. Then $v(f) \leq \text{capacity}(A, B)$.

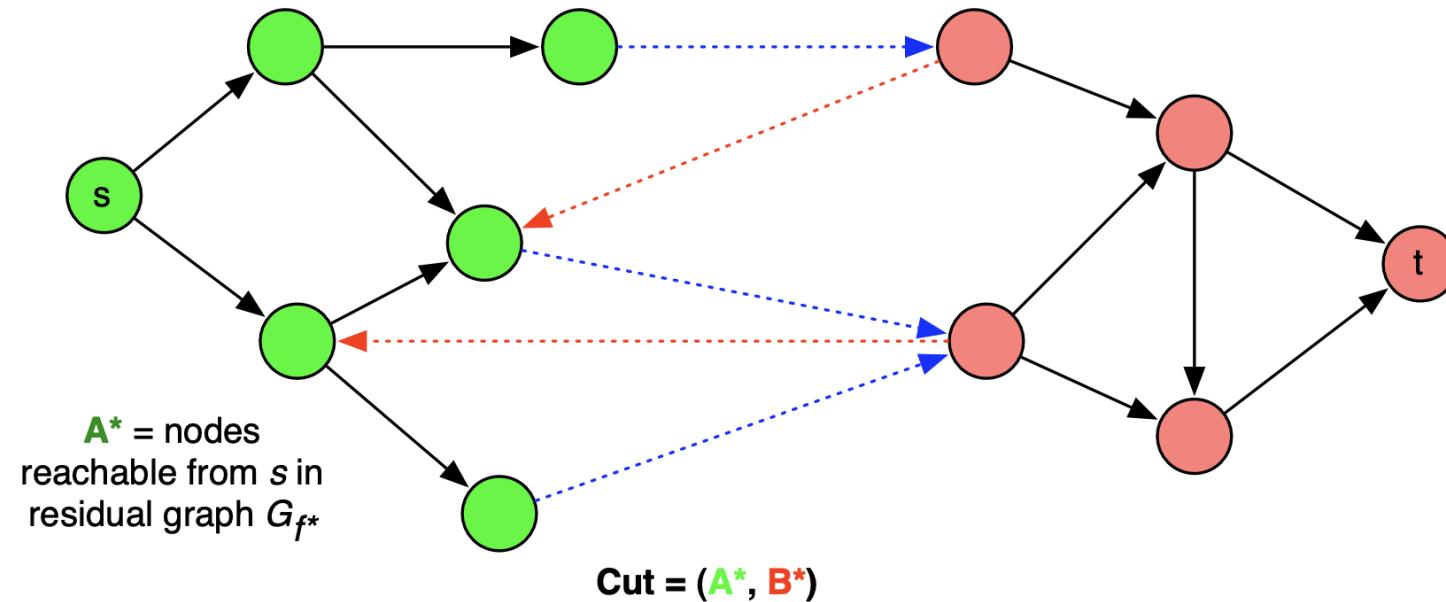
- Says that any cut is bigger than any flow.
- Therefore, cuts constrain flows. The minimum capacity cut constrains the maximum flow the most.
- In fact, the capacity of the minimum cut always equals the maximum flow value.

Max-Flow = Min-Cut



Max-Flow = Min-Cut

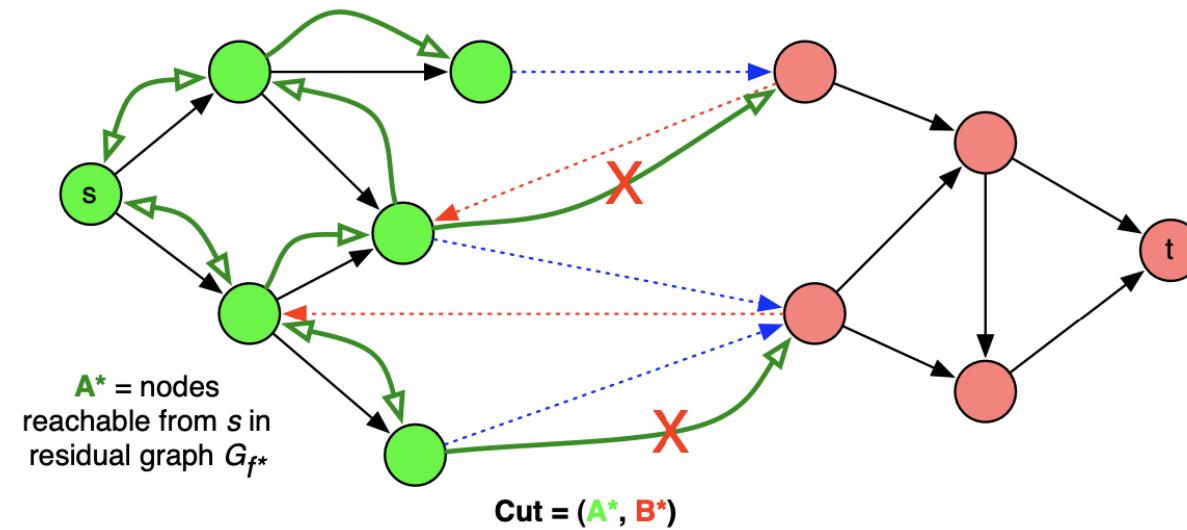
- Let f^* be the flow returned by our algorithm. Look at G_{f^*} but define a cut in G :



Blue edges must be saturated.

Red edges must have 0 flow $\implies v(f^*) = \text{capacity}(A^*, B^*)$.

Max-Flow = Min-Cut



- (A^*, B^*) is an s - t cut because there is no path from s to t in the residual graph G_{f^*} .
- Edges (u,v) from A^* to B^* must be saturated — otherwise there would be a forward edge (u,v) in G_{f^*} and v would be part of A^* .
- Edges (v,u) from B^* to A^* must be empty — otherwise there would be a back edge (u,v) in G_{f^*} and v would be part of A^* .

Max-Flow = Min-Cut

Therefore,

- $v(f^*) = \text{capacity}(A^*, B^*)$.
- No flow can have value bigger than $\text{capacity}(A^*, B^*)$.
- So, f^* must be an maximum flow.
- And (A^*, B^*) has to be a minimum-capacity cut.

Theorem (Max-flow = Min-cut): The value of the maximum flow in any flow graph is equal to the capacity of the minimum cut.

Reference

- L. R. Ford and D. R. Fulkerson. Flows in Networks. Princeton University Press, 1962.

Finding the Min-Capacity Cut

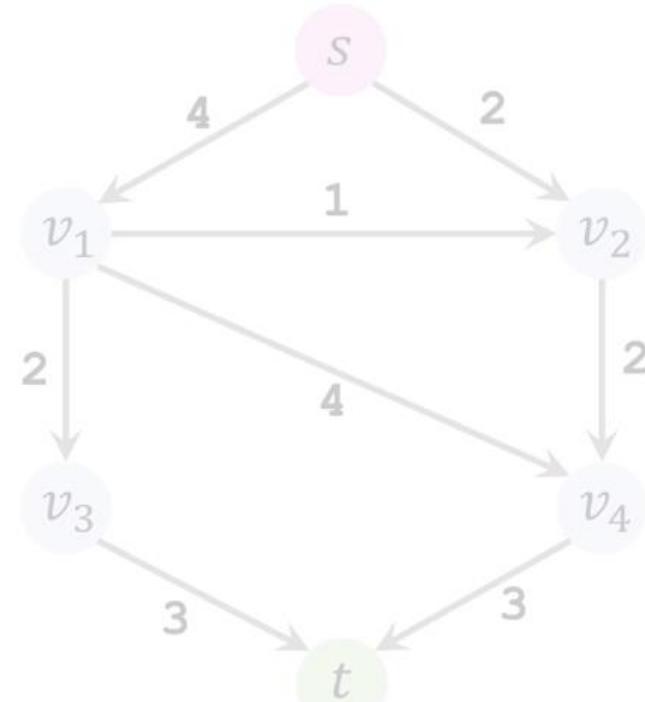
Our proof that maximum flow = minimum cut can be used to actually find the minimum capacity cut:

- Find the maximum flow f^* .
- Construct the residual graph G_{f^*} for f^* .

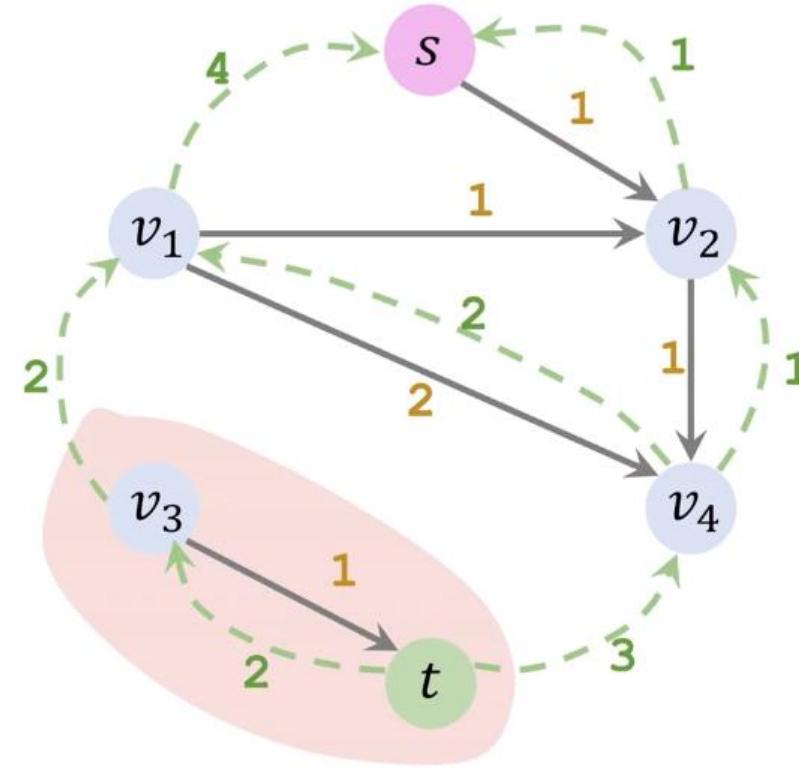
- Do a BFS to find the nodes reachable from s in G_{f^*} . Let the set of these nodes be called A^* .
- Let B^* be all other nodes.
- Return (A^*, B^*) as the minimum capacity cut.

Example

Original Graph



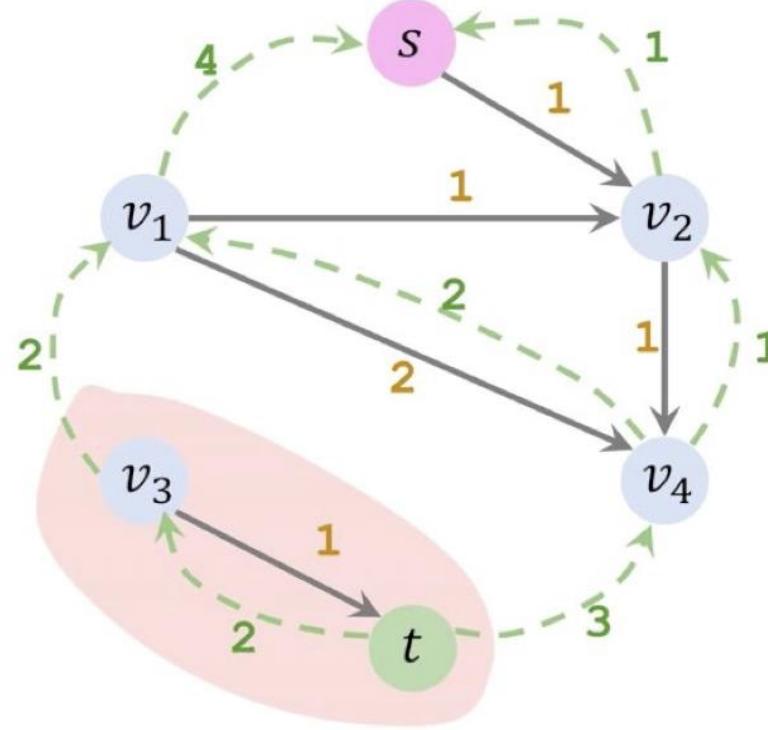
Final Residual Graph



Example

- Find vertices reachable from S
- Subset S = {s, v₁, v₂, v₄}.
- The remaining vertices are t and v₃
- Subset T = {t, v₃}
- Q: What is the min-cut then?

Office on the web Frame
Residual Graph



Applications

- The Maximum Flow Minimum Cut Theorem is a pivotal concept in network flow theory. This theorem has a wide range of applications across various fields. Here are some specific examples of its applications:
 - Water Resources Management
 - Resource Allocation
 - Transportation and Logistics
 - Communications Networks
 - Economic Decision-Making
 - Image Processing
 - Medical Field
 - Social Network Analysis

Transportation and Logistics

- Water Distribution Networks: Ensuring the maximum water flow from a source to various users.
- Energy Distribution: Optimizing the maximum electricity transmission from power stations to consumers in power networks.

Communications Networks

- Bandwidth Allocation: Ensuring the maximum data transmission rate from source to destination on the internet.
- Routing Algorithms: Choosing paths to maximize network throughput in data packet transmission.
- Production Process Optimization: Optimizing the maximum flow from raw materials to finished goods while minimizing production bottlenecks.

Summary

- Ford-Fulkerson algorithm can find max flow in $O(mC)$ time.
 - Algorithm idea: Send flow along some path with capacity left, possibly “erasing” some flow we’ve already sent. Use residual graph to keep track of remaining capacities and flow we’ve already sent.
- We can eliminate C to get a true polynomial algorithm by using BFS to find our augmenting paths.
- All cuts have capacity \geq the value of all flows.
- Know the flow is maximum because its value equals the capacity of some cut.

The End