



香港科技大学(广州)
THE HONG KONG
UNIVERSITY OF SCIENCE AND
TECHNOLOGY (GUANGZHOU)

Design and Analysis of Algorithms

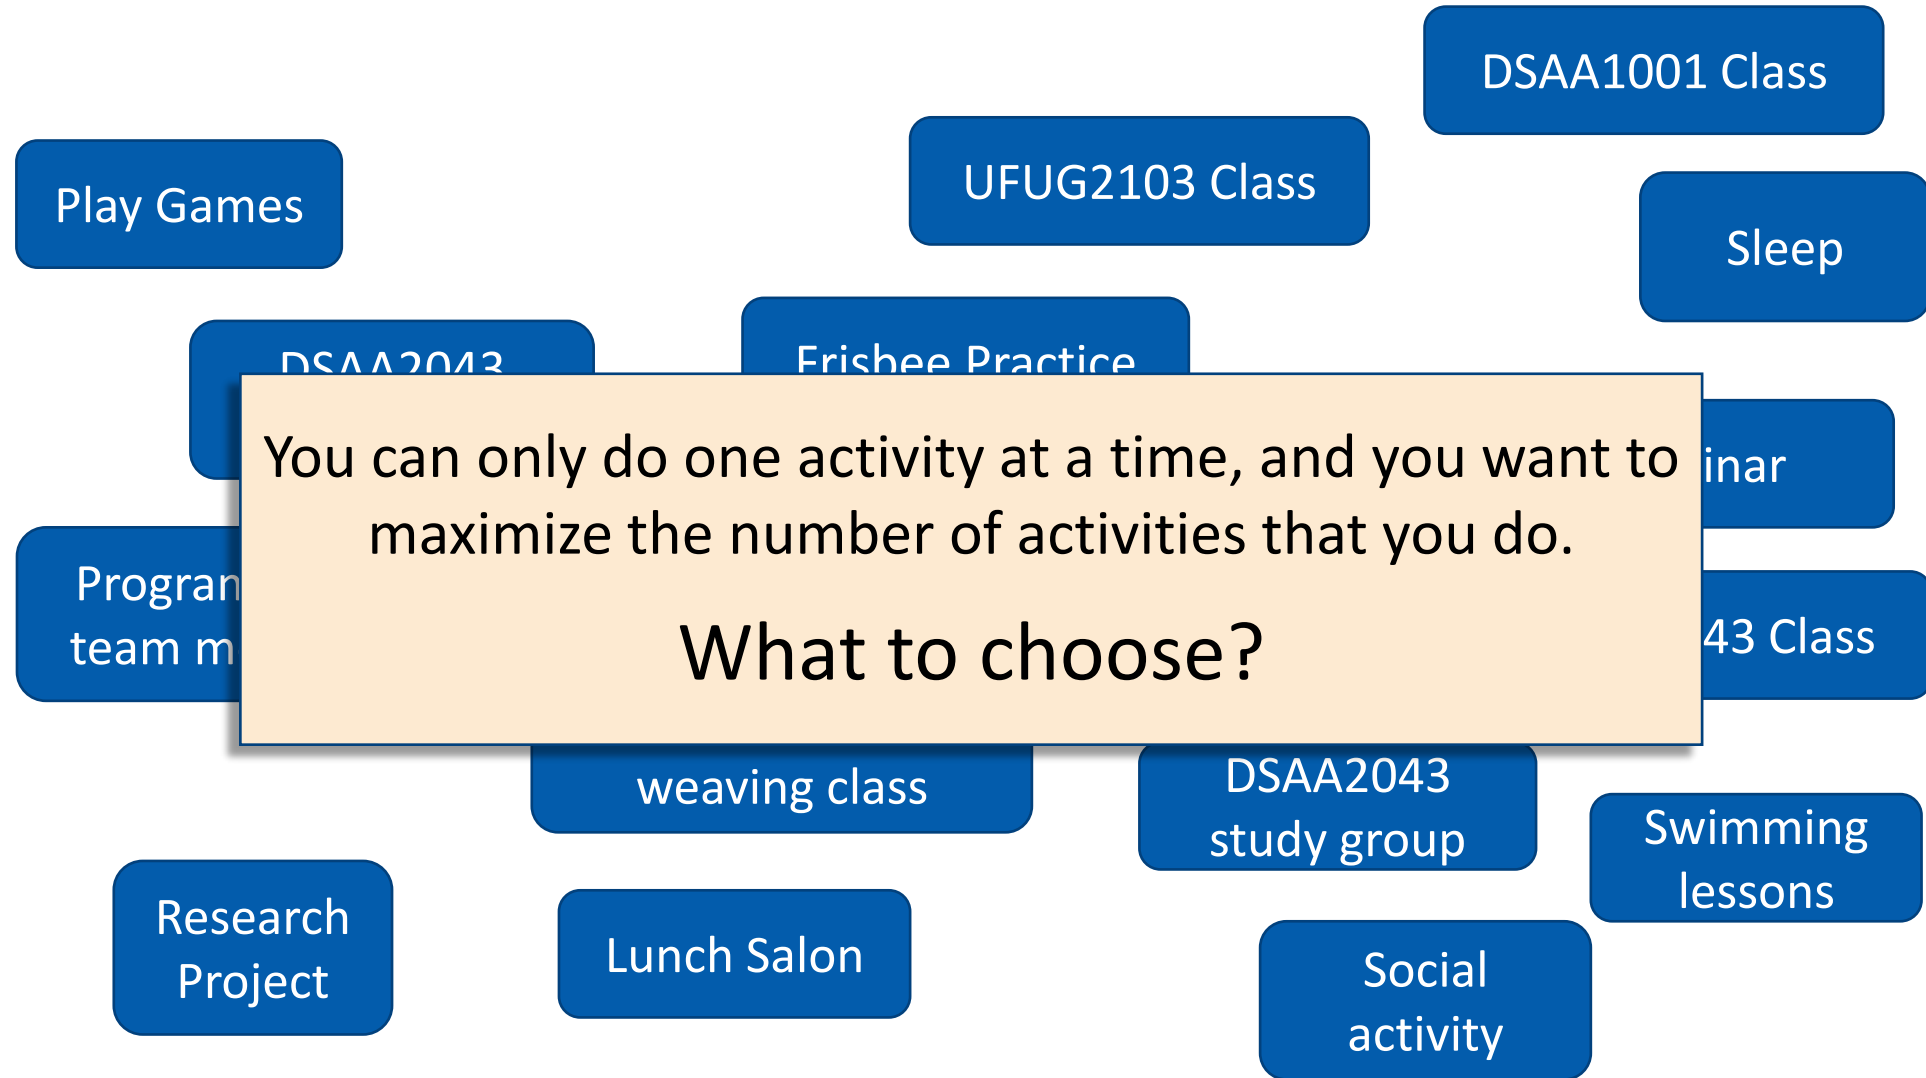
Jing Tang | DSAA 2043 Fall 2024

Greedy Algorithms

One example of a **greedy algorithm** that **does not work**:
Knapsack

Three examples of **greedy algorithms** that **do work**:
Activity Selection
Job Scheduling
Minimum Spanning Tree

Example where greedy works

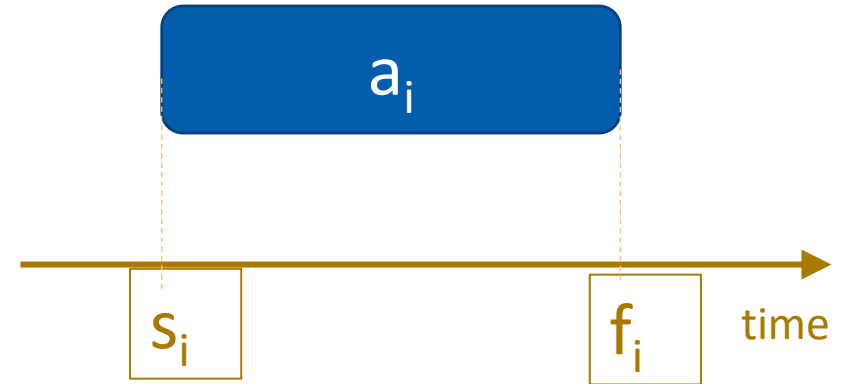


- Input:

- Activities a_1, a_2, \dots, a_n
- Start times s_1, s_2, \dots, s_n
- Finish times f_1, f_2, \dots, f_n

- Output:


- A way to maximize the number of activities you can do today.



In what order should you greedily add activities?

In what order?

- Shortest job first?



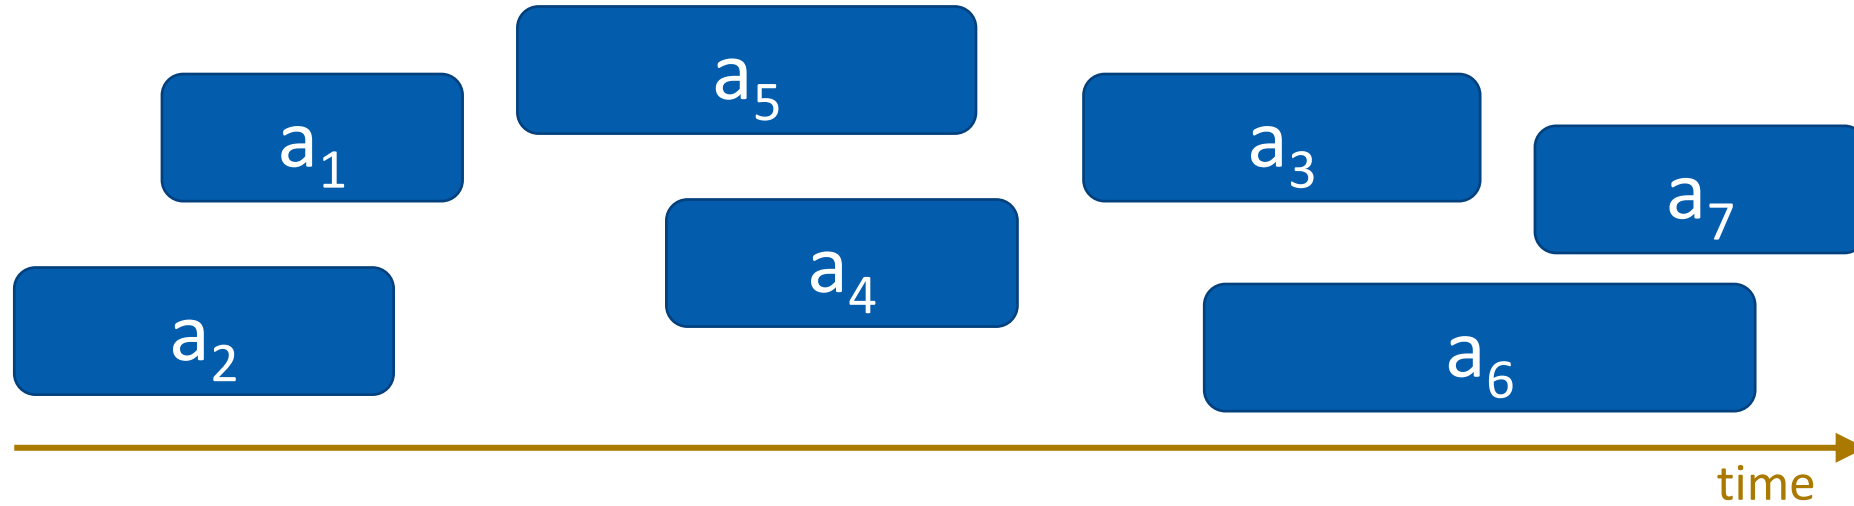
- Earliest start time?



- Earliest finish time?

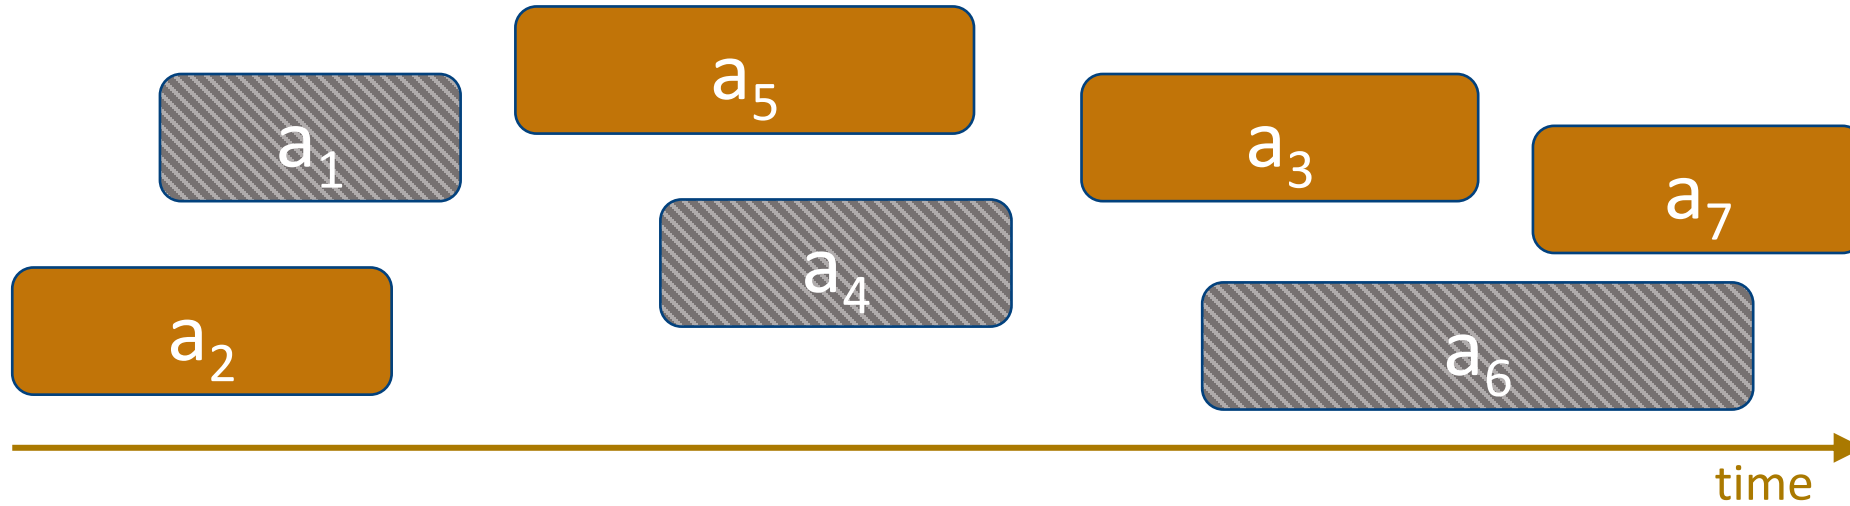


Greedy Algorithm



- Pick activity you can add with the smallest finish time.
- Repeat.

Greedy Algorithm



- Pick activity you can add with the smallest finish time.
- Repeat.

- Running time:
 - $O(n)$ if the activities are already sorted by finish time.
 - Otherwise, $O(n \log(n))$ if you have to sort them first.

Why does it work?

- **We never rule out an optimal solution**
- At the end of the algorithm, we've got some solution.
- So it must be optimal.

A common strategy for proving the correctness of greedy algorithms:

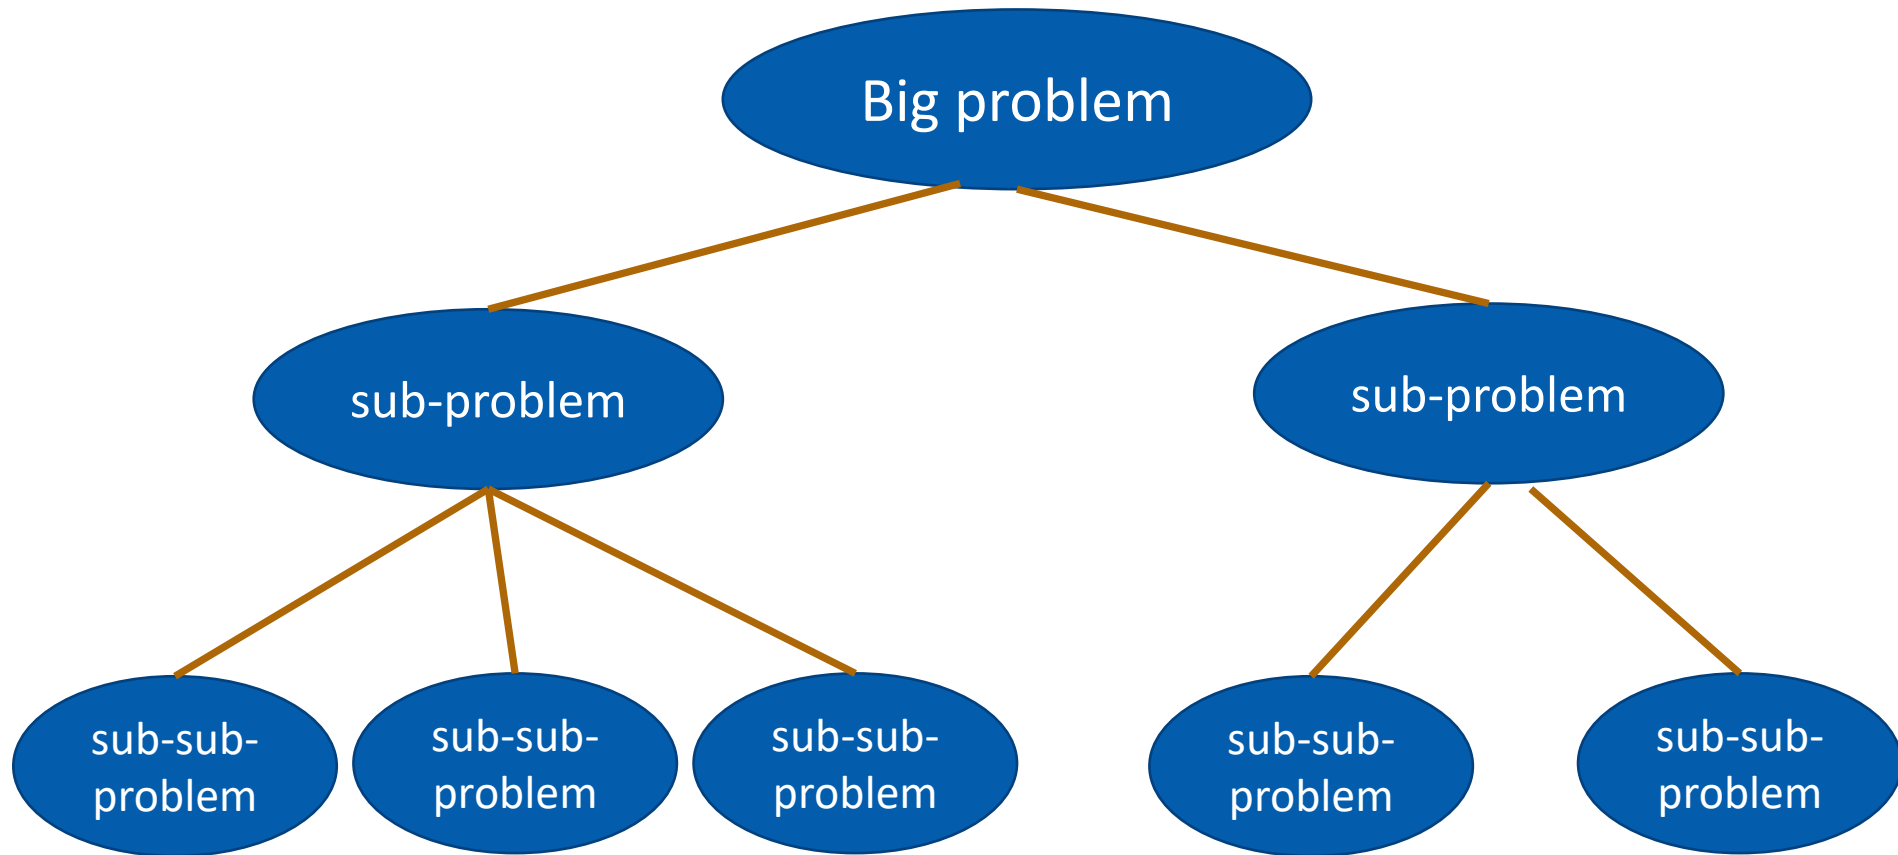
- Make a **series of choices**.
- Show that, at each step, our choice **won't rule out an optimal solution** at the end of the day.
- After we've made all our choices, we haven't ruled out an optimal solution, **so we must have found one**.

- Inductive Hypothesis:
 - After greedy choice t , you haven't ruled out success.
- Base case:
 - Success is possible before you make any choices.
- Inductive step:
 - If you haven't ruled out success after choice t , then you won't rule out success after choice $t+1$.
- Conclusion:
 - If you reach the end of the algorithm and haven't ruled out success then you must have succeeded.

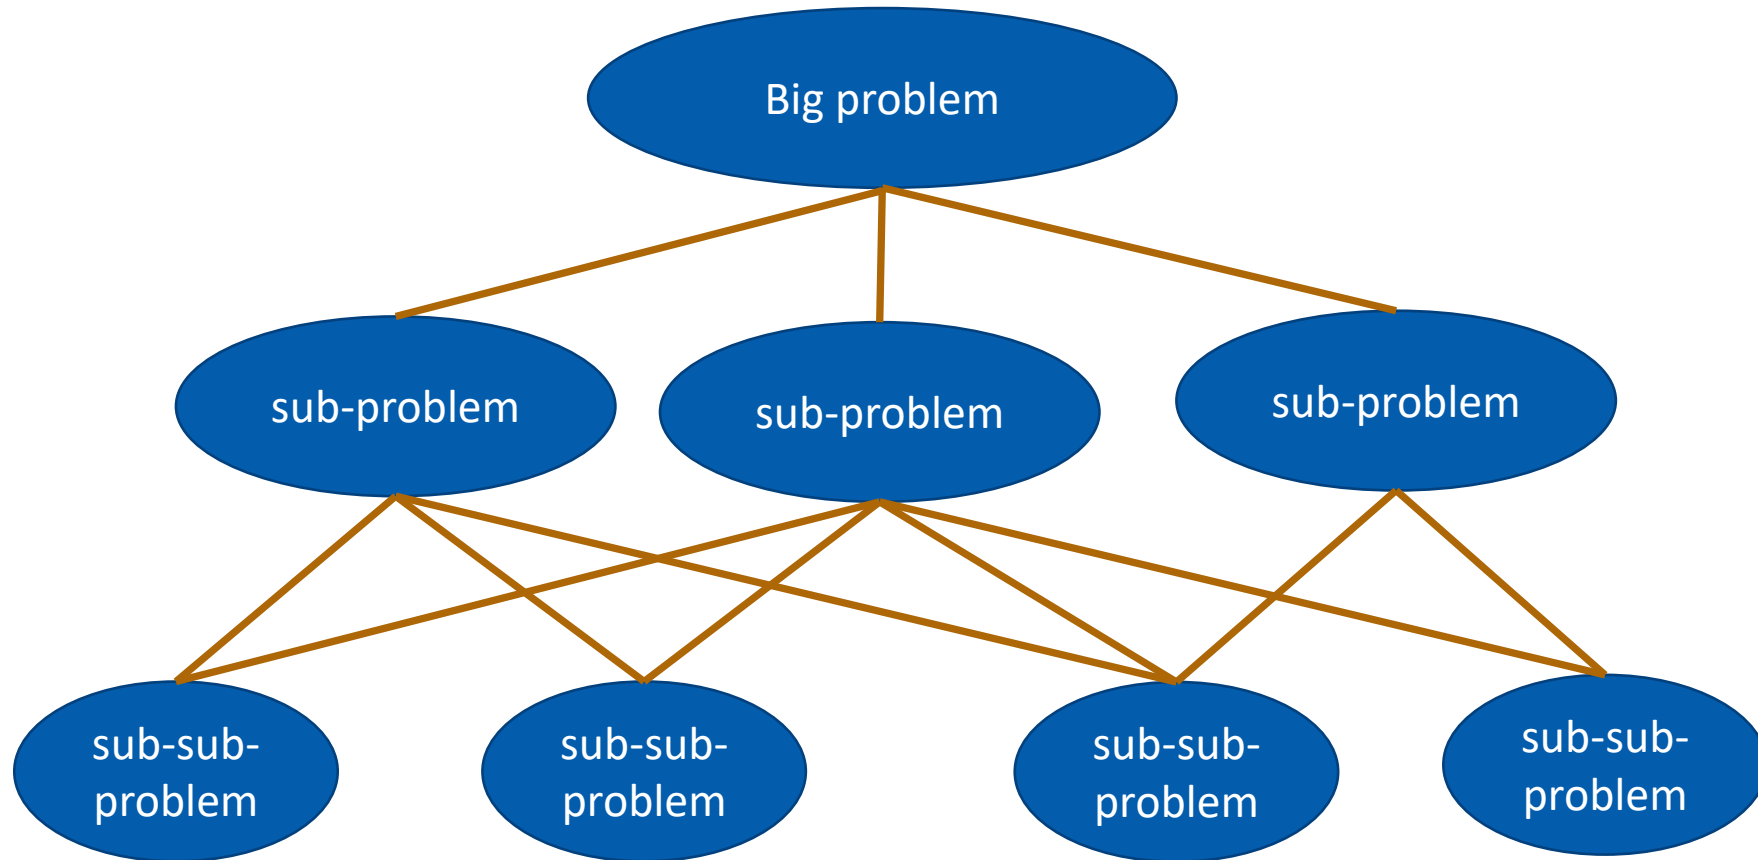
A common strategy for showing we don't rule out the optimal solution:

- Suppose that you're on track to make an optimal solution T^* .
 - E.g., after you've picked activity i , you're still on track.
- Suppose that T^* *disagrees* with your next greedy choice.
 - E.g., it *doesn't* involve activity k .
- Manipulate T^* in order to make a solution T that's not worse but that *agrees* with your greedy choice.
 - E.g., swap whatever activity T^* did pick next with activity k .


- Divide-and-conquer:



- Dynamic Programming:



- Greedy algorithms:



- Not only is there **optimal sub-structure**:
 - optimal solutions to a problem are made up from optimal solutions of sub-problems
- but each problem **depends on only one sub-problem**.

Another Example: Scheduling

DSAA2043 HW

Personal hygiene

Math HW

Administrative stuff for student club

Econ HW

Do laundry

Sports


Practice musical instrument

Read lecture notes

Have a social life

Sleep

- n tasks
- Task i takes t_i hours
- For every hour that passes until task i is done, pay c_i



- DSAA2043 HW, then Sleep: costs $10 \cdot 2 + (10 + 8) \cdot 3 = 74$ units
- Sleep, then DSAA2043 HW: costs $8 \cdot 3 + (10 + 8) \cdot 2 = 60$ units

- Seems amenable to a greedy algorithm:

Take the best job first

Then solve this problem



Take the best job first

Then solve this problem



Take the best job first

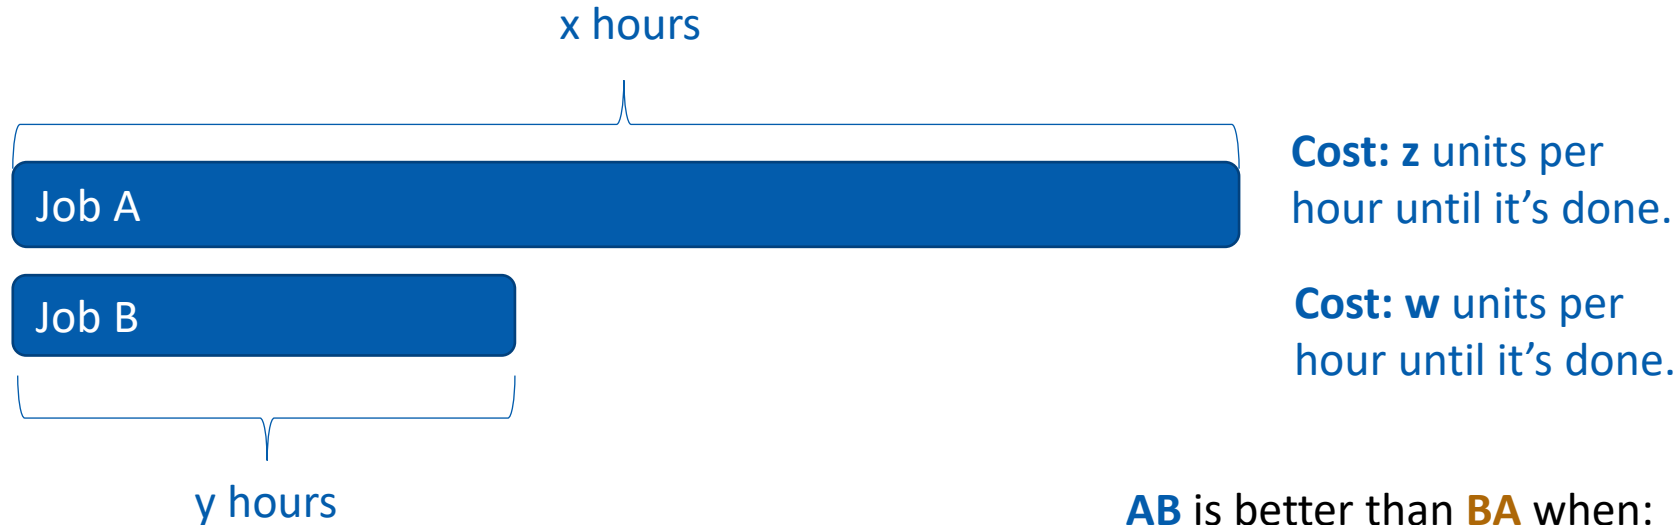
Then solve this problem



(That one's easy 😊)

What does “best” mean?

- Of these two jobs, which should we do first?



- Cost(**A then B**) = $x \cdot z + (x + y) \cdot w$
- Cost(**B then A**) = $y \cdot w + (x + y) \cdot z$

AB is better than **BA** when:

$$xz + (x + y)w \leq yw + (x + y)z$$

$$xz + xw + yw \leq yw + xz + yz$$

$$wx \leq yz$$


$$\frac{w}{y} \leq \frac{z}{x}$$

- Choose the job with the biggest $\frac{\text{cost of delay}}{\text{time it takes}}$ ratio.

Minimum Spanning Trees

Minimum Spanning Trees

- Say we have an undirected weighted graph




A **spanning tree** is a **tree** that connects all of the vertices.

A **tree** is a connected graph with no cycles!

Minimum Spanning Trees

- Say we have an undirected weighted graph

The **cost** of a spanning tree is the sum of the weights on the edges.




This is a spanning tree with cost 67.

A **spanning tree** is a **tree** that connects all of the vertices.

Minimum Spanning Trees

- Say we have an undirected weighted graph




minimum

of minimum cost

A **spanning tree** is a **tree** that connects all of the vertices.

Why MSTs?


- Network design
 - Connecting cities with roads/electricity/telephone/...
- Cluster analysis
 - E.g., genetic distance
- Image processing
 - E.g., image segmentation
- Useful primitive
 - For other graph algs



How to find an MST


Idea:

Start growing a tree, greedily add the shortest edge we can to grow the tree.



Idea:

Start growing a tree, greedily add the shortest edge we can to grow the tree.



We've discovered Prim's algorithm!


- `slowPrim(G = (V,E), starting vertex s)`:
 - `MST = {}`
 - `verticesVisited = { s }`
 - **while** `|verticesVisited| < |V|`:
 - find the lightest edge `{x,v}` in `E` so that:
 - `x` is in `verticesVisited`
 - `v` is not in `verticesVisited`
 - add `{x,v}` to `MST`
 - add `v` to `verticesVisited`
 - **return** `MST`

Naively, the running time is $O(nm)$:

- For each of $\leq n-1$ iterations of the **while** loop:
 - Go through all the edges.


Efficient Implementation

- Each vertex keeps:
 - the **(single-edge) distance** from itself to the **growing spanning tree**
 - **how to get there.**



Efficient Implementation

- Each vertex keeps:
 - the **(single-edge) distance** from itself to the **growing spanning tree**
 - **how to get there.**
- Choose the closest vertex, add it.




I'm 7 away.
C is the closest.

I can't get to the
tree in one edge


Efficient Implementation

- Each vertex keeps:
 - the **(single-edge) distance** from itself to the **growing spanning tree**
 - **how to get there.**
- Choose the closest vertex, add it.
- Update stored info.



Efficient Implementation

Every vertex has a key and a parent




Can't reach x yet

x is "active"

Can reach x



$k[x]$ is the distance of x from the growing tree




$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.

Until all the vertices are **reached**:

- Activate the **unreached** vertex u with the **smallest key**.
- **for each** of u's unreached neighbors v:
 - $k[v] = \min(k[v], \text{weight}(u,v))$
 - if $k[v]$ updated, $p[v] = u$

Efficient Implementation

Every vertex has a key and a parent




Can't reach x yet

x is "active"

Can reach x



$k[x]$ is the distance of x from the growing tree



$p[b] = a$, meaning that a was the vertex that $k[b]$ comes from.

Until all the vertices are **reached**:

- Activate the **unreached** vertex u with the **smallest key**.
- **for each** of u's unreached neighbors v:
 - $k[v] = \min(k[v], \text{weight}(u,v))$
 - if $k[v]$ updated, $p[v] = u$


- Very similar to Dijkstra's algorithm!
- **Differences:**
 1. Keep track of $p[v]$ in order to return a tree at the end
 - But Dijkstra's can do that too, that's not a big difference.
 2. Instead of $d[v]$ which we update by
 - $d[v] = \min(d[v], d[u] + w(u,v))$we keep $k[v]$ which we update by
 - $k[v] = \min(k[v], w(u,v))$

Thing 2 is the
main difference.

Kruskal's Algorithm

what if we just always take the cheapest edge?
whether or not it's connected to what we have so far?


That won't
cause a cycle



- **slowKruskal**($G = (V, E)$):
 - Sort the edges in E by non-decreasing weight.
 - $MST = \{\}$
 - **for** e in E (in sorted order):
 - **if** adding e to MST won't cause a cycle:
 - add e to MST .
 - **return** MST
- m iterations through this loop
- How do we check this?


Kruskal's Algorithm

At each step of Kruskal's, we are maintaining a forest.



Kruskal's Algorithm

At each step of Kruskal's, we are maintaining a forest.
When we add an edge, we merge two trees:




Union-find data structure

Implementation – Lab 13

- Used for storing collections of sets
- Supports:
 - **makeSet(u)**: create a set {u}
 - **find(u)**: return the set that u is in
 - **union(u,v)**: merge the set that u is in with the set that v is in.

```
makeSet(x)  
makeSet(y)  
makeSet(z)
```

```
union(x, y)  
find(x)
```



- **kruskal**($G = (V, E)$):
 - Sort E by weight in non-decreasing order
 - $MST = \{\}$ // initialize an empty tree
 - **for** v in V :
 - **makeSet**(v) // put each vertex in its own tree in the forest
 - **for** (u, v) in E : // go through the edges in sorted order
 - **if** **find**(u) \neq **find**(v): // if u and v are not in the same tree
 - add (u, v) to MST
 - **union**(u, v) // merge u 's tree with v 's tree
 - **return** MST

Running time

- Sorting the edges takes $O(m \log(n))$
 - In practice, if the weights are small integers we can use radixSort and take time $O(m)$
- For the rest:
 - n calls to **makeSet**
 - put each vertex in its own set
 - $2m$ calls to **find**
 - for each edge, **find** its endpoints
 - $n-1$ calls to **union**
 - we will never add more than $n-1$ edges to the tree,
 - so we will never call **union** more than $n-1$ times.
- Total running time: **$O(m \log(n))$**

Complexity Classes

- Definition: The class P consists of all decision problems that are solvable in polynomial time
- Definition: The class NP consists of all decision problems such that, for each yes-input, there exists a certificate that can be verified in polynomial time.
 - NP stands for “**Nondeterministic Polynomial time**”.
- $P = NP$?

The End